DEPLOYING HIGHLY AVAILABLE, SCALABLE AND SECURE WEBSITE ON AWS

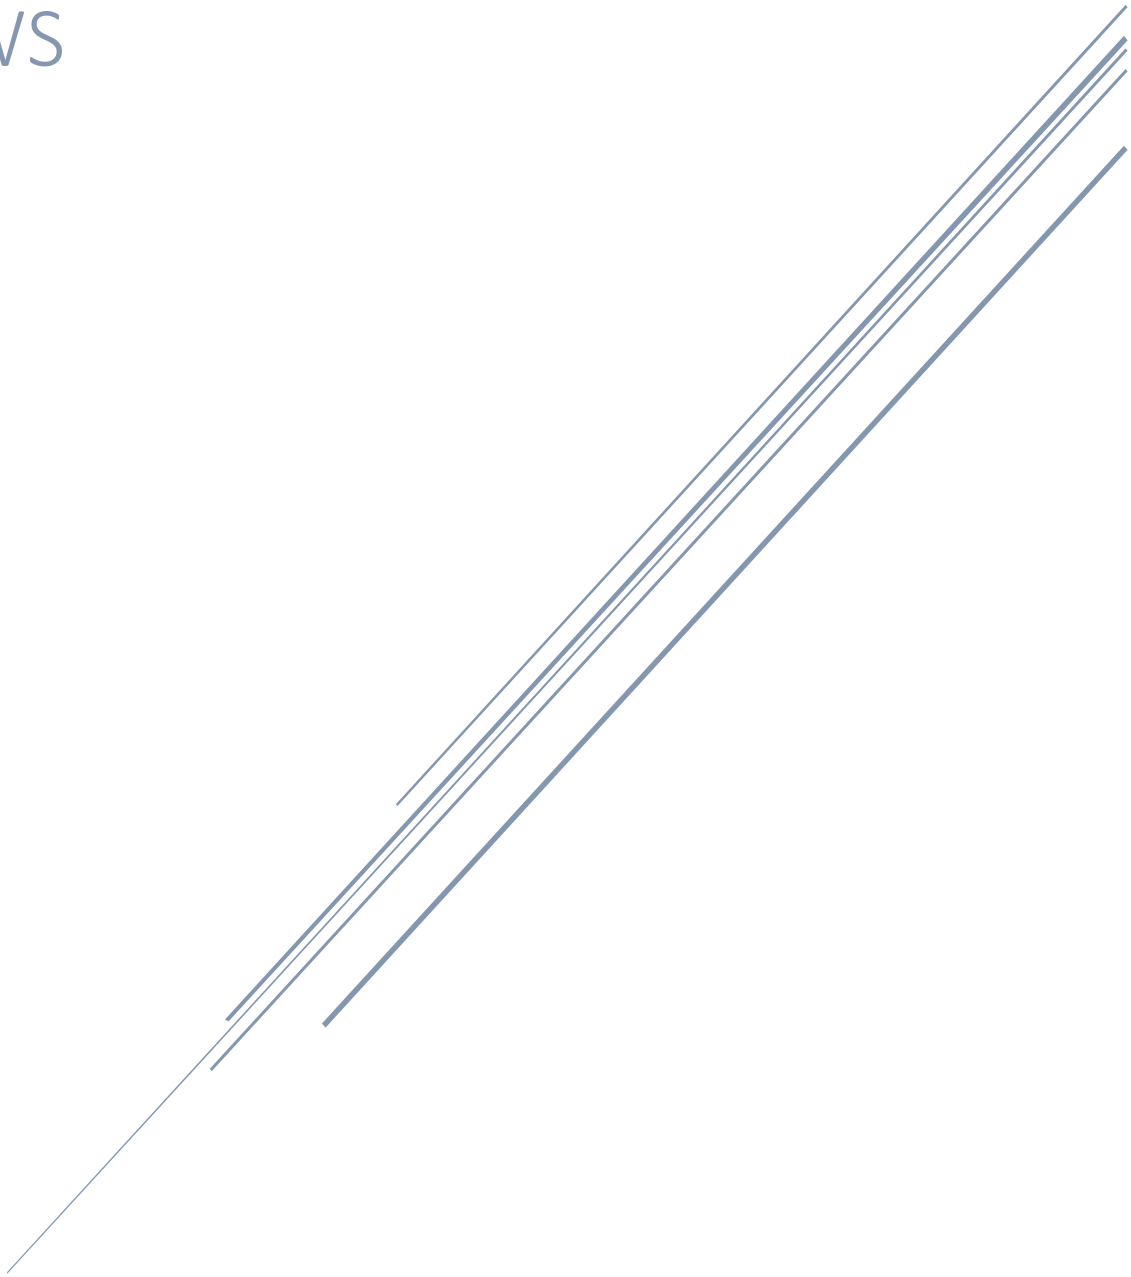

VINAY SHARMA

Table of Contents

Deploying Highly Available, Scalable and Secure Website on AWS

Diagram: High Availability and Secure Web Application on AWS

This diagram illustrates the architecture of my web application deployed on AWS, featuring load balancing, multi-AZ deployment, a custom VPC with private subnets, and enhanced security measures including a custom domain and SSL certificate.

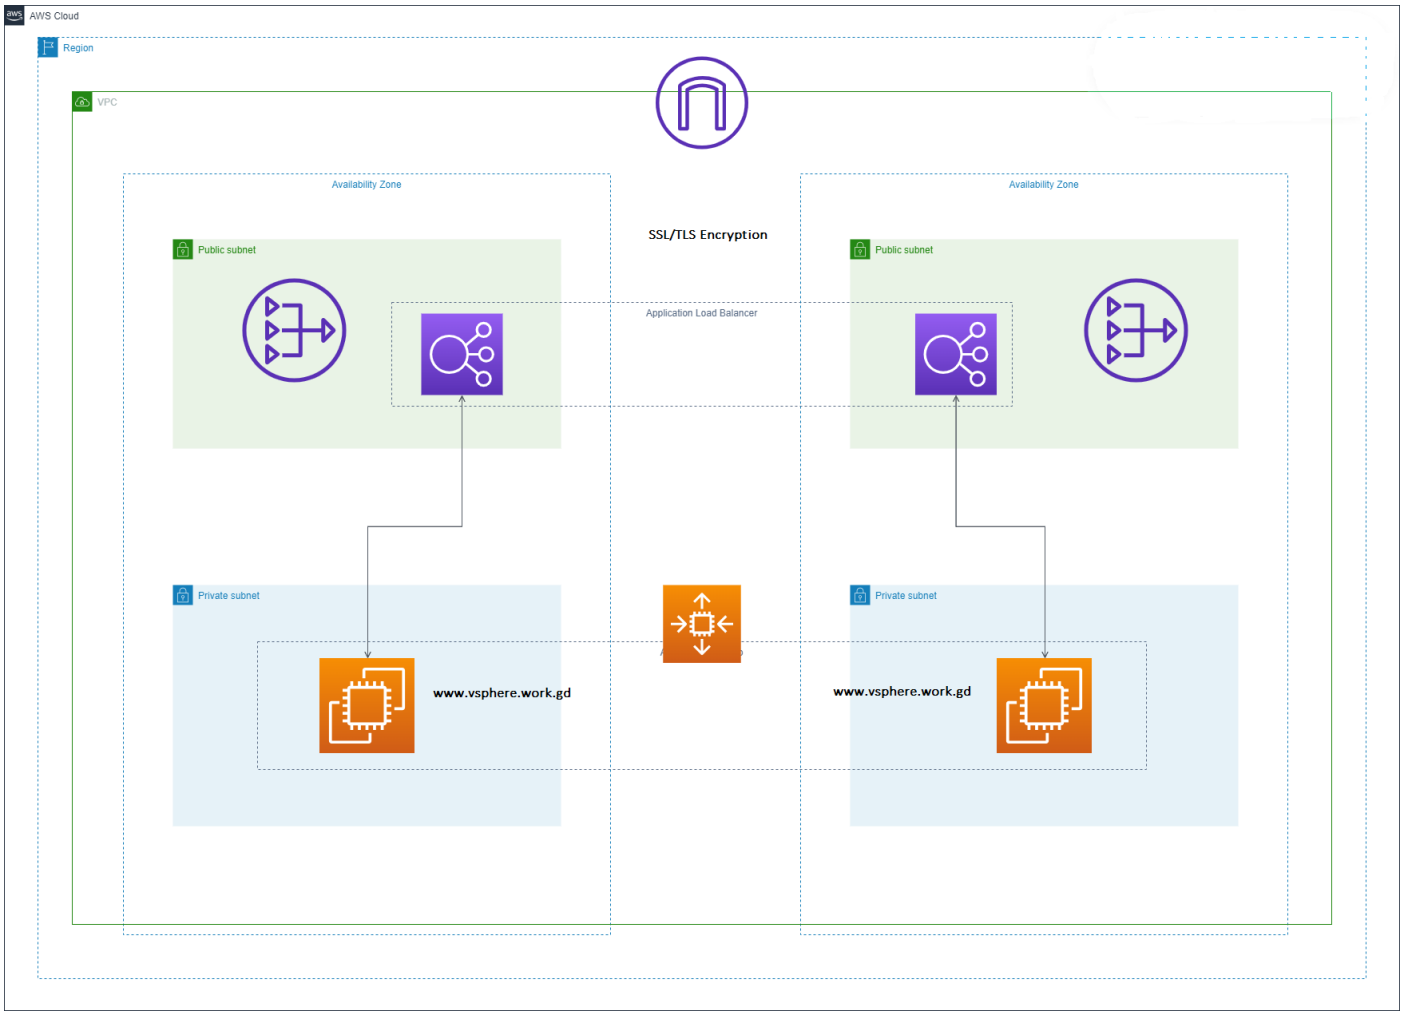

Project Description: High Availability, Scalability and Security of Web Application on AWS

In this project, I have deployed a sample website on AWS, implementing a robust
architecture to ensure high availability, security, and scalability. The key components and configurations are as follows:

1. **Load Balancing:** To evenly distribute incoming traffic and enhance fault tolerance, I have employed an AWS Elastic Load Balancer (ELB). This ensures that user requests are efficiently handled by the available servers.
2. **Multi-AZ Deployment:** The web application is hosted across multiple Availability Zones (AZs) within a region. This setup guarantees high availability and resilience, as it mitigates the impact of an AZ failure.
3. **Custom VPC:** A custom Virtual Private Cloud (VPC) has been created to provide network isolation and enhanced security. This allows for fine-grained control over the network environment.
4. **Private Subnets:** The servers hosting the website are placed in private subnets, enhancing security by restricting direct access from the internet. These subnets are configured to route traffic through the load balancer, which is the only publicly accessible endpoint.

5. **Security Groups and NACLs:** Security groups and Network Access Control Lists (NACLs) are configured to control inbound and outbound traffic at both the instance and subnet level, ensuring that only legitimate traffic reaches the servers.

6. **Domain and SSL Certificate:** A custom domain has been purchased for the website, providing a professional and memorable address for users.
   Additionally, an SSL certificate has been generated and implemented to secure the data transmission between the users and the application, ensuring that all communications are encrypted and trustworthy.

7. **Access through Load Balancer:** The web application is only accessible through the load balancer's public DNS, providing a single entry point and further securing the backend servers.

This architecture not only ensures high availability and fault tolerance but also secures the application by minimizing direct exposure to the internet and controlling access strictly through the load balancer. The addition of a custom domain and SSL certificate further enhances the user experience by providing a secure and reliable connection.

In this project, I got a chance to work on the following technologies:

1. AWS EC2

2. AWS Auto Scaling Group (ASG)

3. AWS Application Load Balancer (ASG)

4. AWS VPC

5. Domain Name System (DNS)

6. Private Subnet

6. SSL Certificates

Practical Implementation:

Let's Create Custom VPC (Virtual Private Cloud):

VPC (Virtual Private Cloud): A VPC is a virtual network that closely resembles a traditional network that you'd operate in your own data center.

A **Custom VPC** in AWS is a user-defined Virtual Private Cloud that allows you to configure and manage your own isolated network within the AWS cloud.

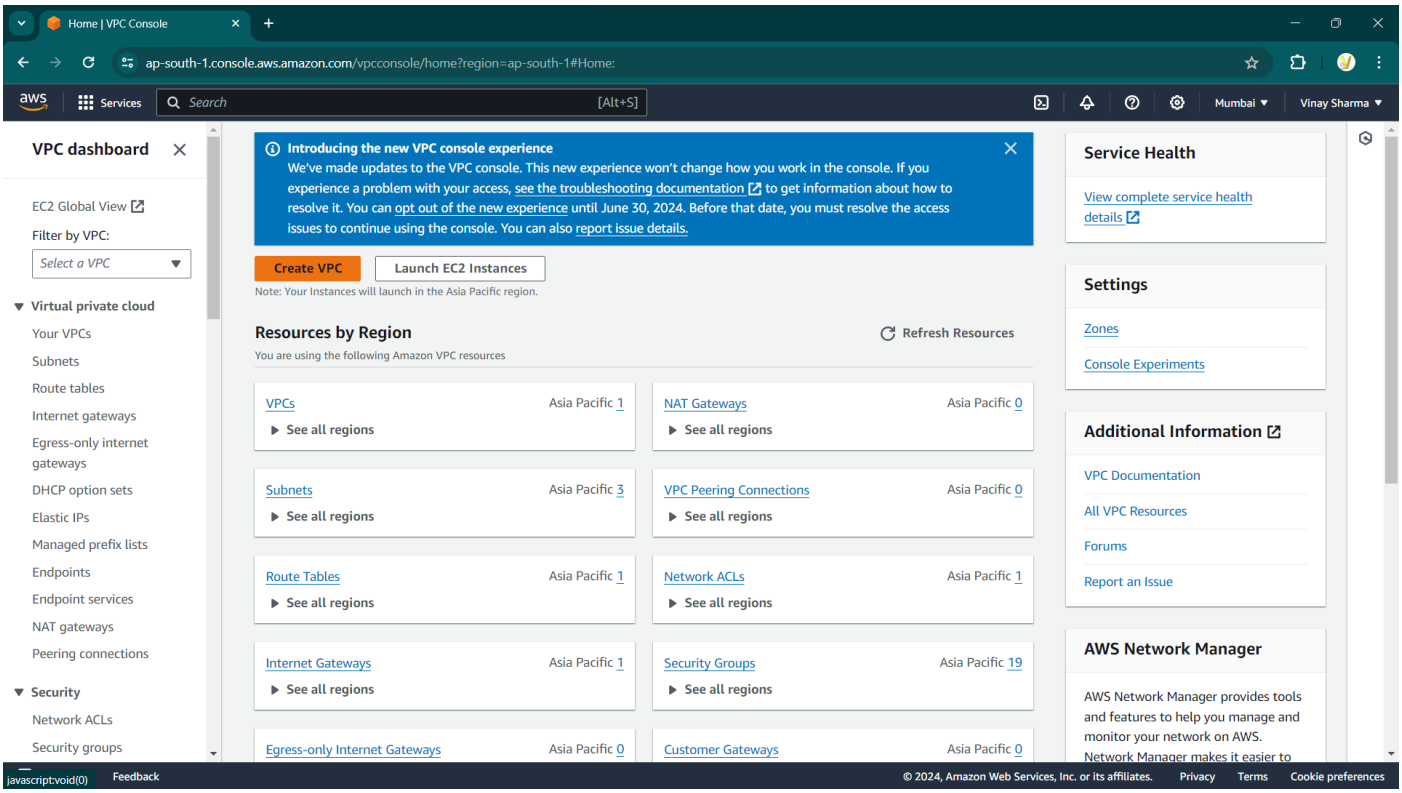

- *Click on create VPC.*

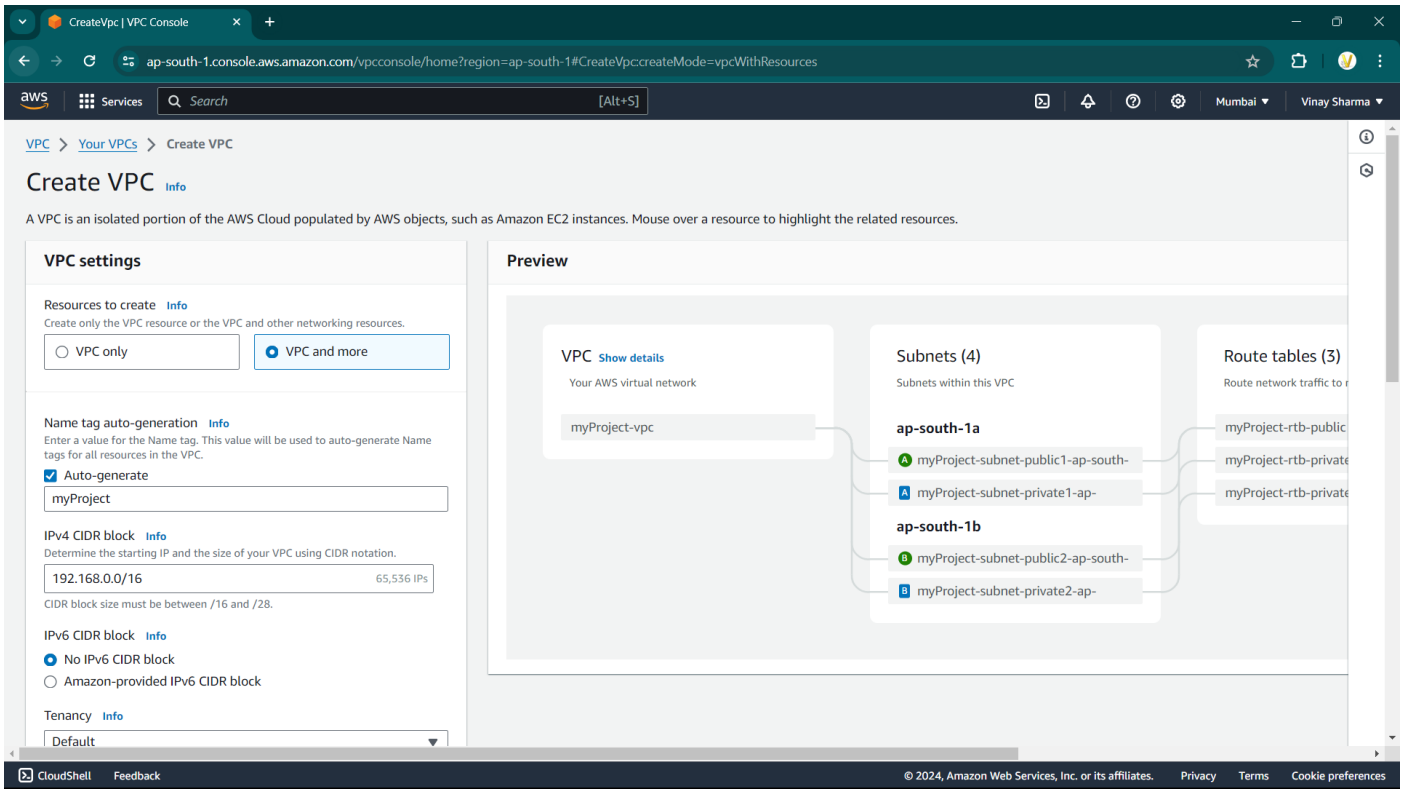

- *Click VPC and More Because it allows you to add VPC and its component together.*

- *IPv4 CIDR Block is nothing but a size on IP address range.*

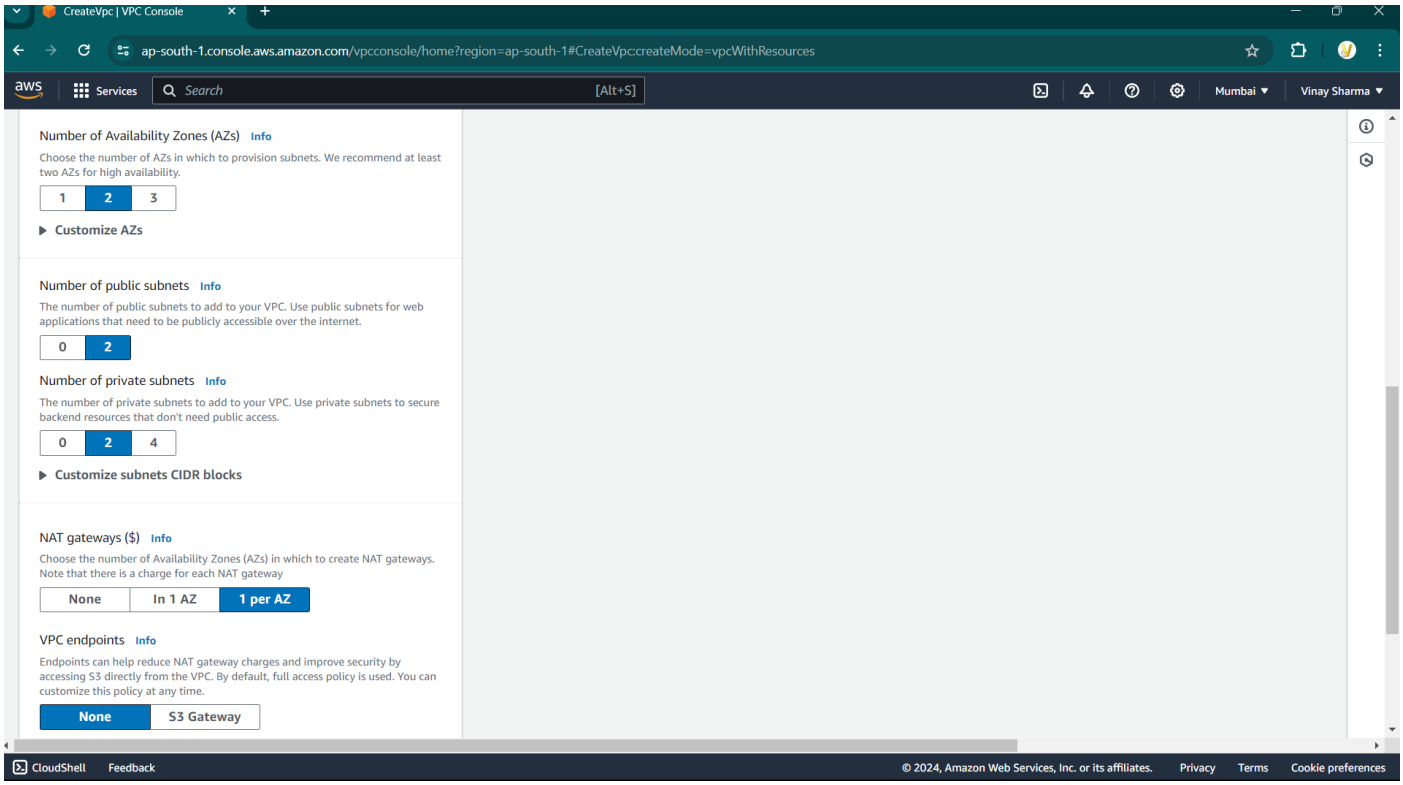

- *I am using Multi-AZ deployment that's why I have choose number of AZ to 2.*
- *In my VPC I want to private and two private subnet in each AZ.*

Public Subnet: has a direct route to an internet gateway. Resources in a public subnet can access the public internet.

Private Subnet: does not have a direct route to an internet gateway.

Internet Gateway: provide two-way public connectivity to applications running in AWS Regions and/or in Local Zones.

NAT Gateway: You can use a NAT gateway so that instances in a private subnet can connect to services outside your VPC but

external services cannot initiate a connection with those instances.

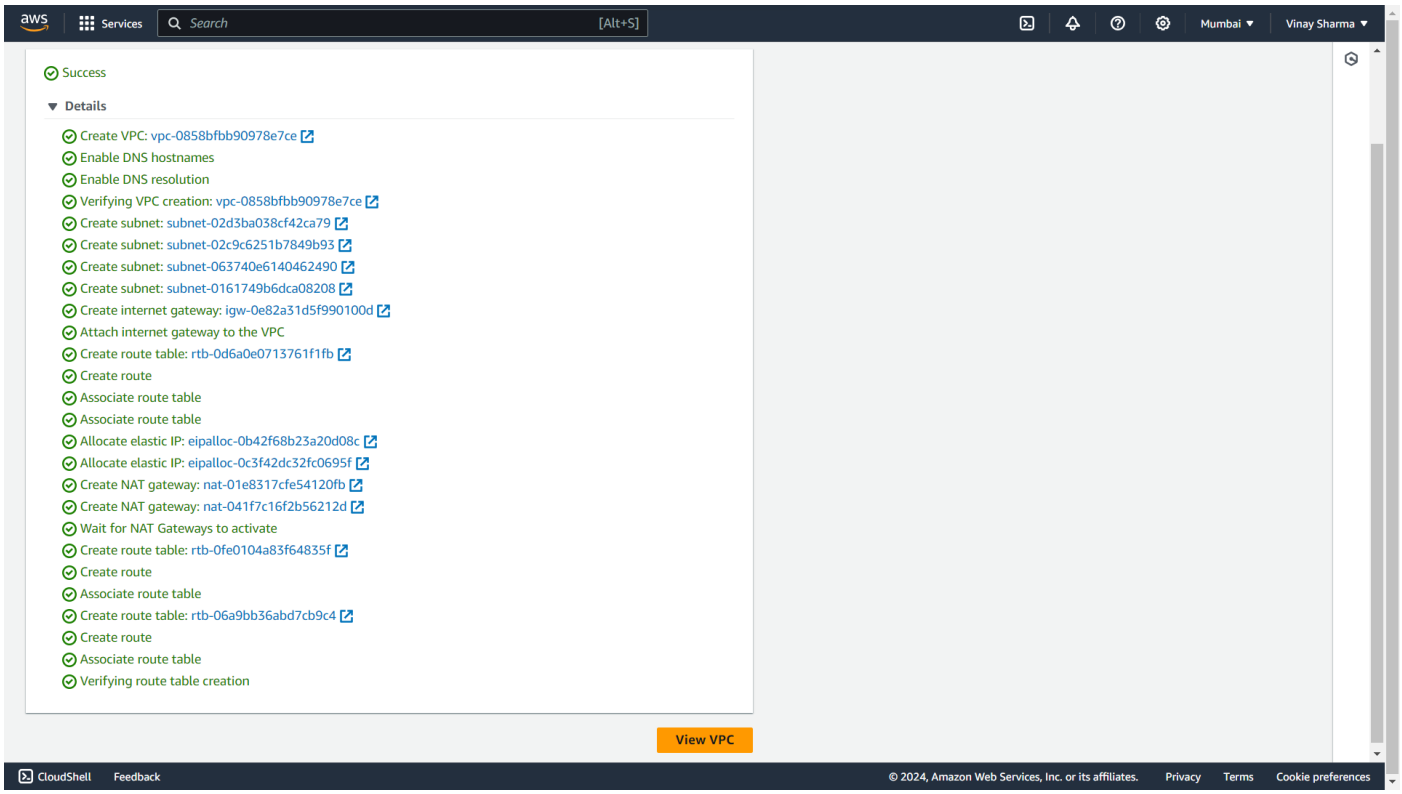

- *Click on create VPC and it will take some time to create VPC.*

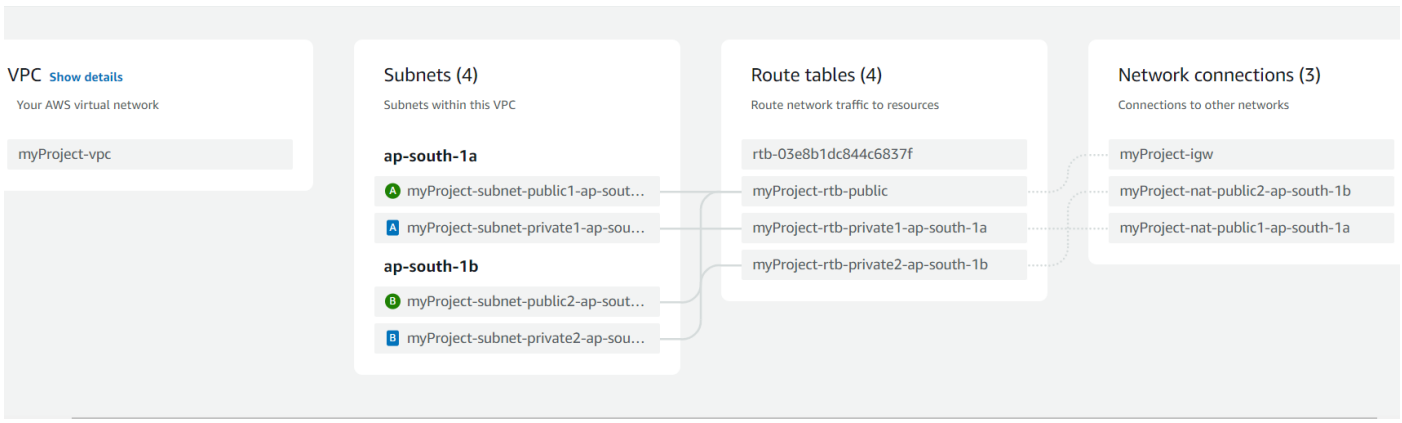

- *Here is our network design.*

Route Table: contains a set of rules, called routes, that determine where network traffic from your subnet or gateway is directed.

Launch Template: holds instance configuration information. It includes the ID of the Amazon Machine Image (AMI), the

instance type, a key pair, security groups, and other parameters used to launch EC2 instances.

Auto Scaling Group: contains a collection of EC2 instances that are treated as a logical grouping for the purposes of automatic scaling and management.

Application Load Balancer: operates at the request level (layer 7), routing traffic to targets (EC2 instances, containers, IP addresses, and Lambda functions) based on the content of the request.

Let's Create Launch Template:

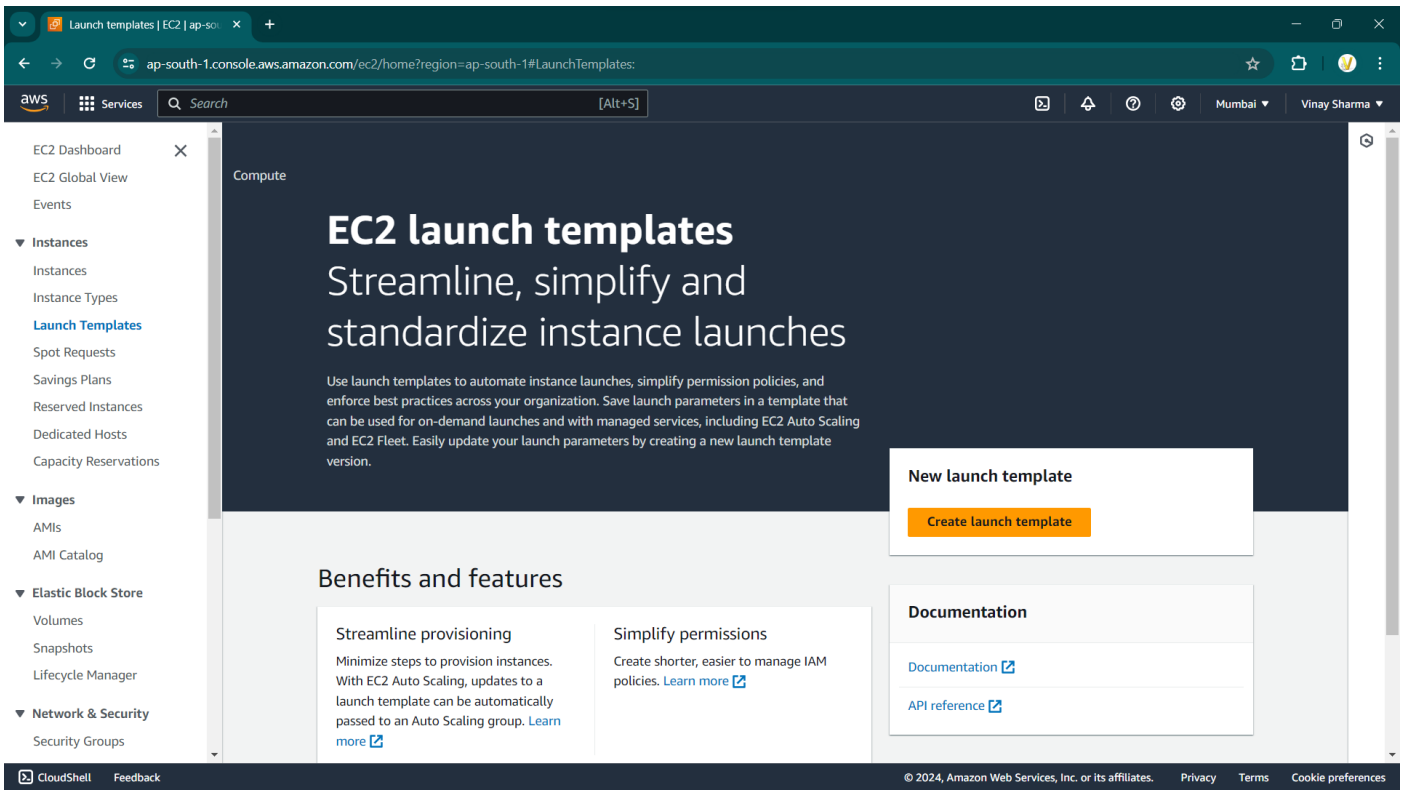

- *Click on create launch template.*

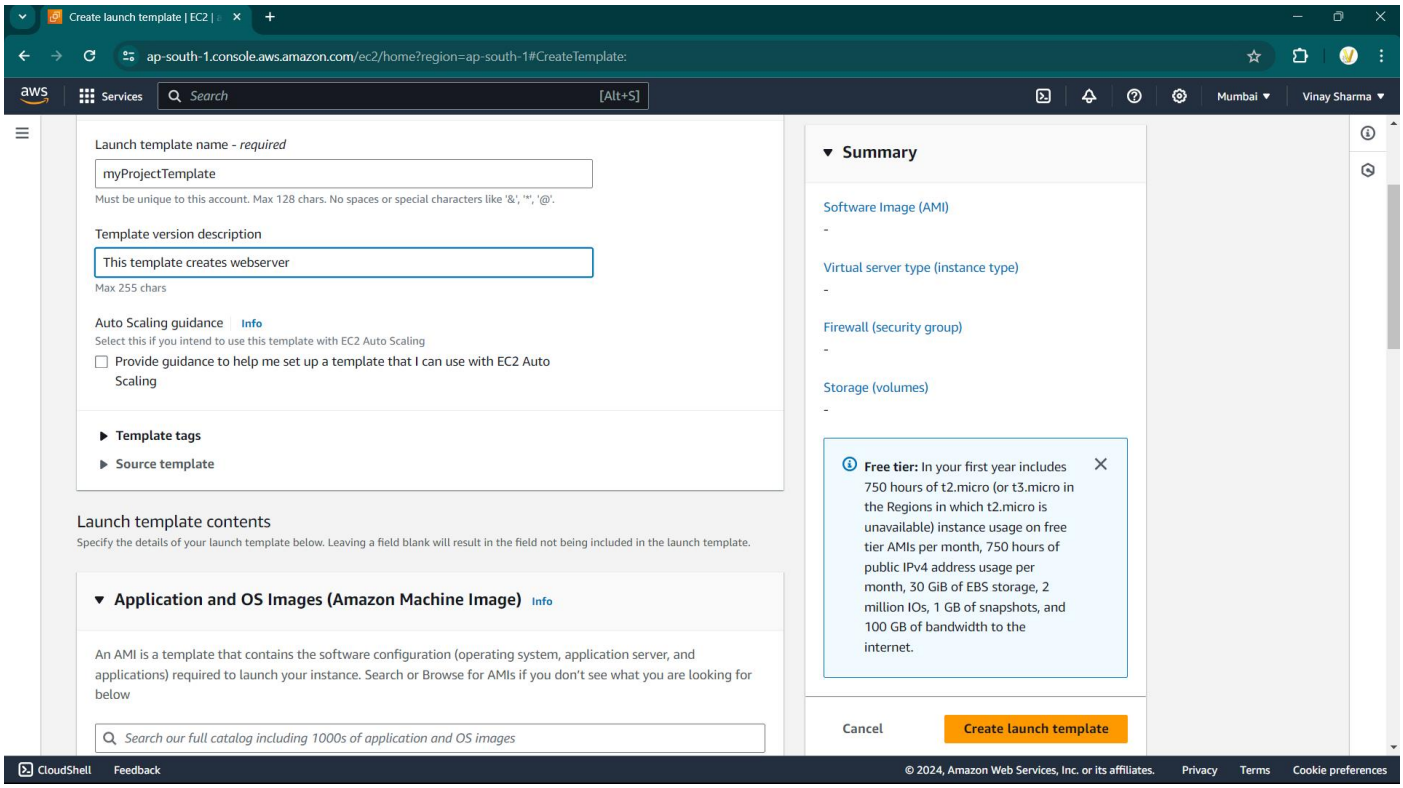

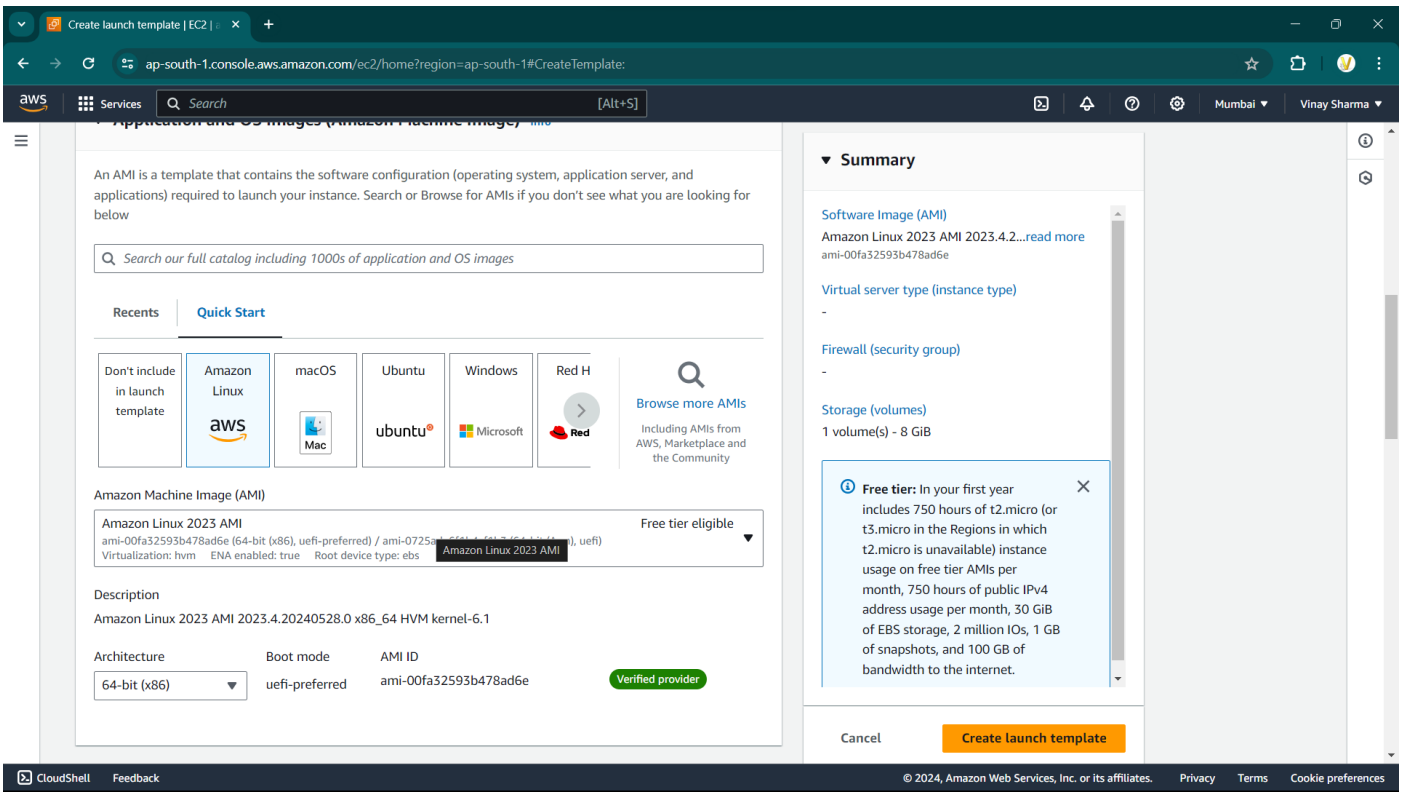

- *Here I am using Amazon Linux 2 as an AMI.*

Amazon Machine Image(AMI): template that contains a software configuration (for example, an operating system, an application server, and applications). From an AMI, you launch an instance.

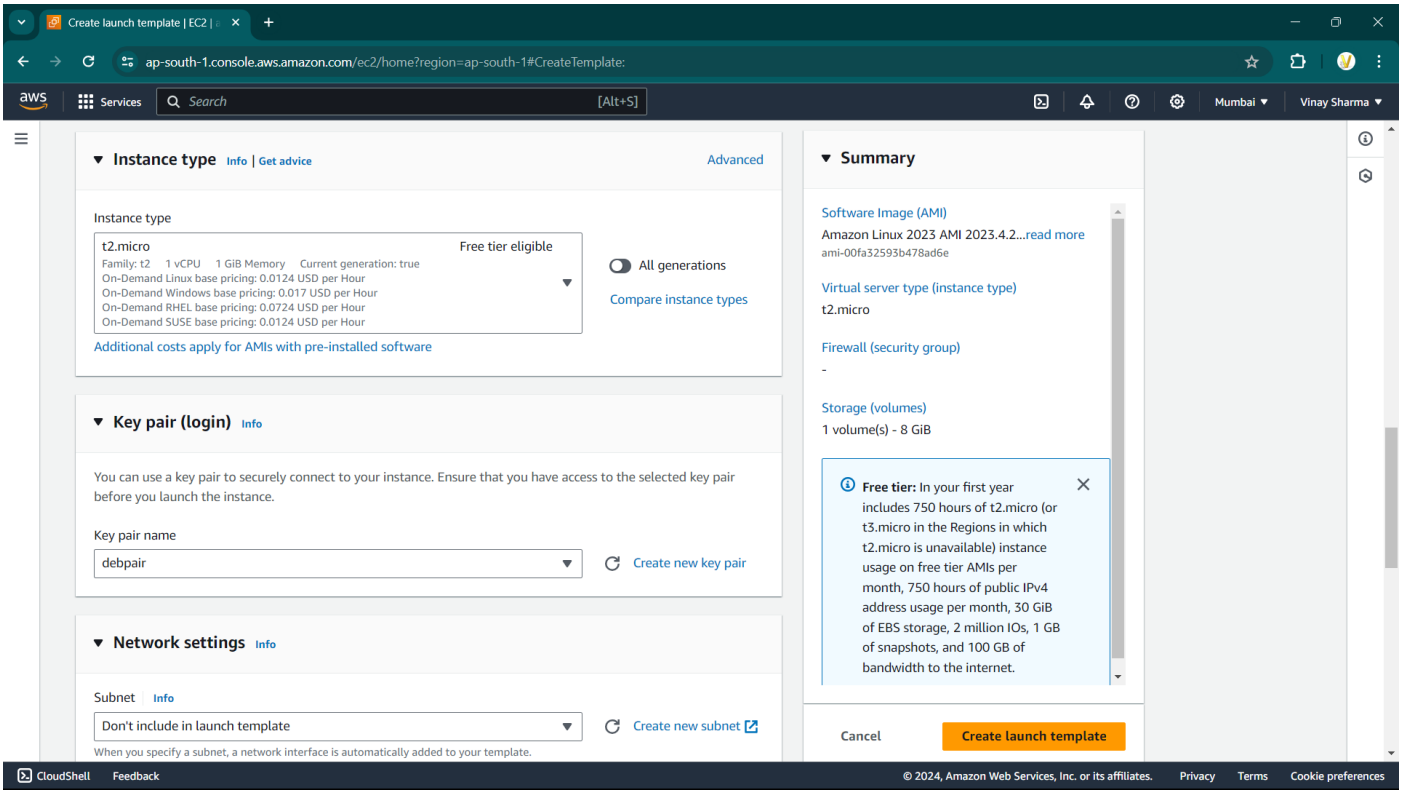

- *Here I am using t2.micro as an instance type which is free one.*

Instance Type: where you specify the hardware of the host computer used for your instance

Key Pair: combination of public and private key that you can use to securely login on your instance.

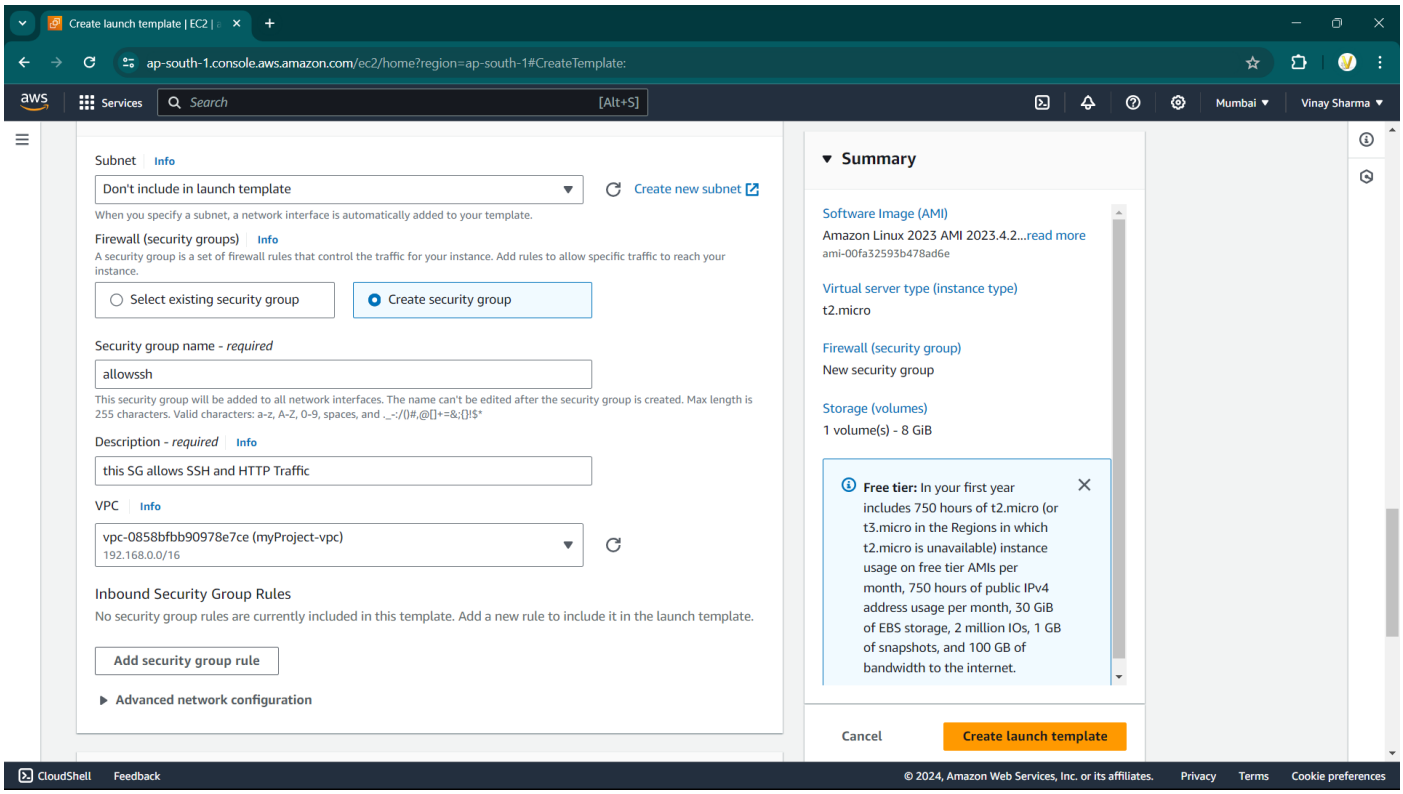

- *Here You have to select the Custom VPC that we have created.*

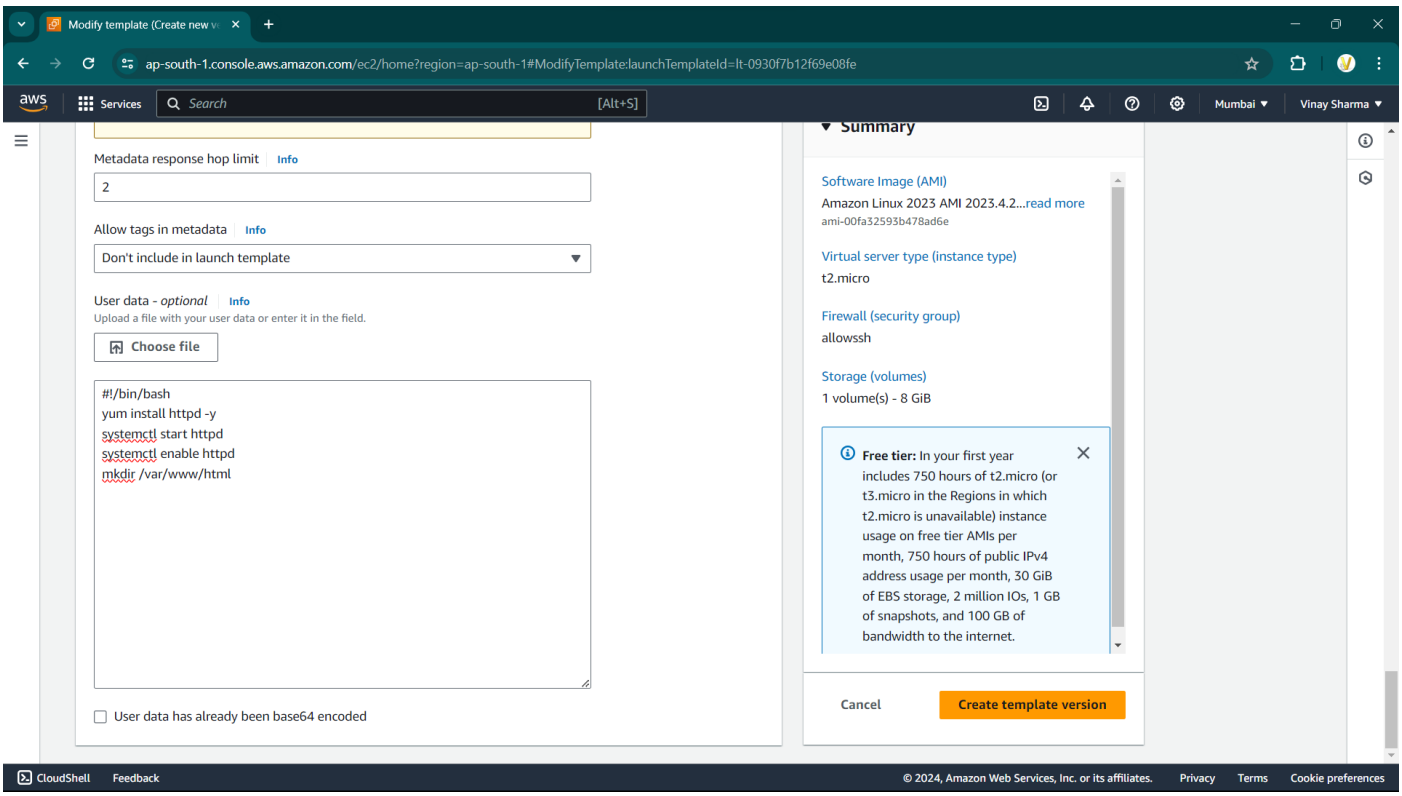

In user data type following command:

```
#!/bin/bash
yum install httpd -y
systemctl start httpd
systemctl enable httpd
mkdir /var/www/html
```

User data: allows user to specify code, task or commands that are going to run at the time of creation of EC2.

Let's Create Auto Scaling Group & Application Load Balancer:

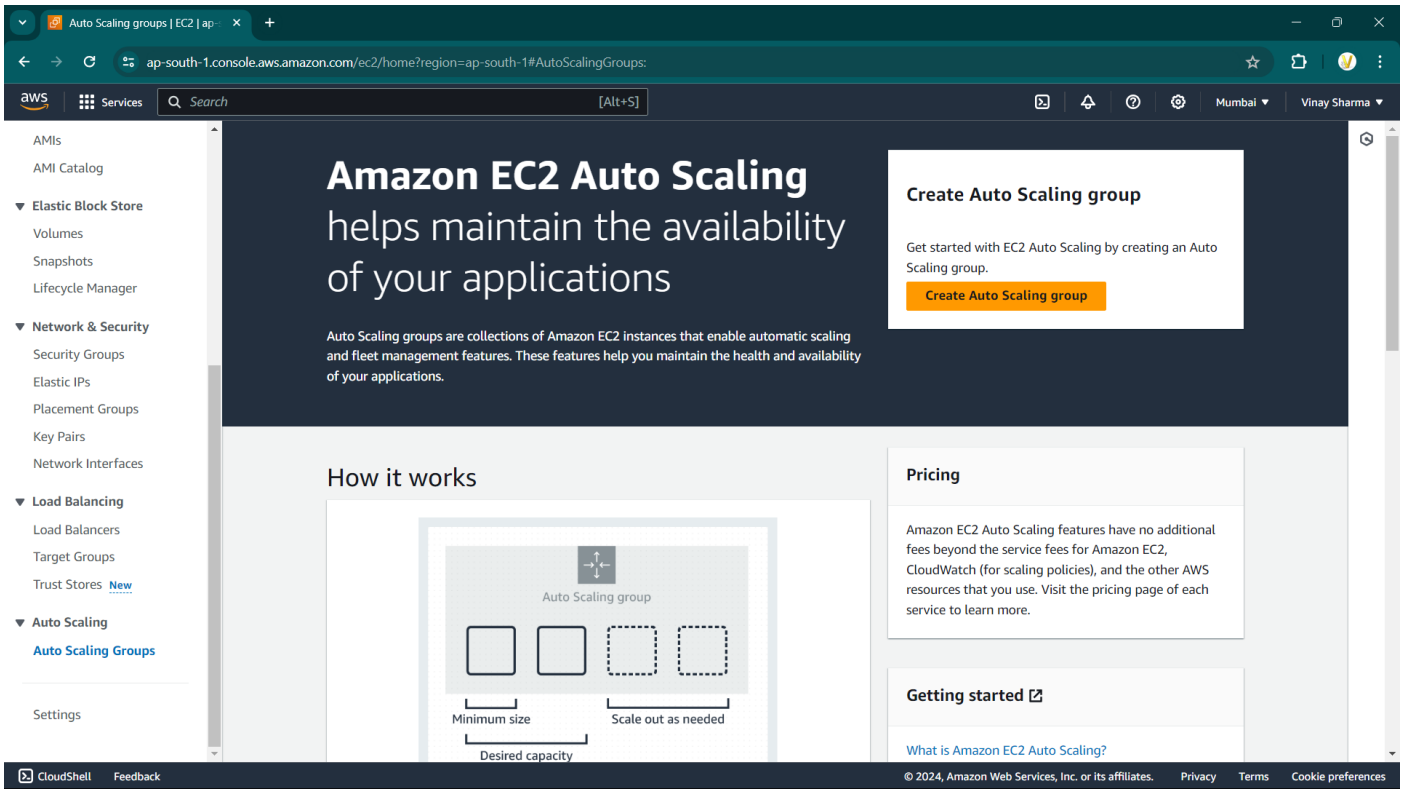

- *Click on create Auto Scaling Group*

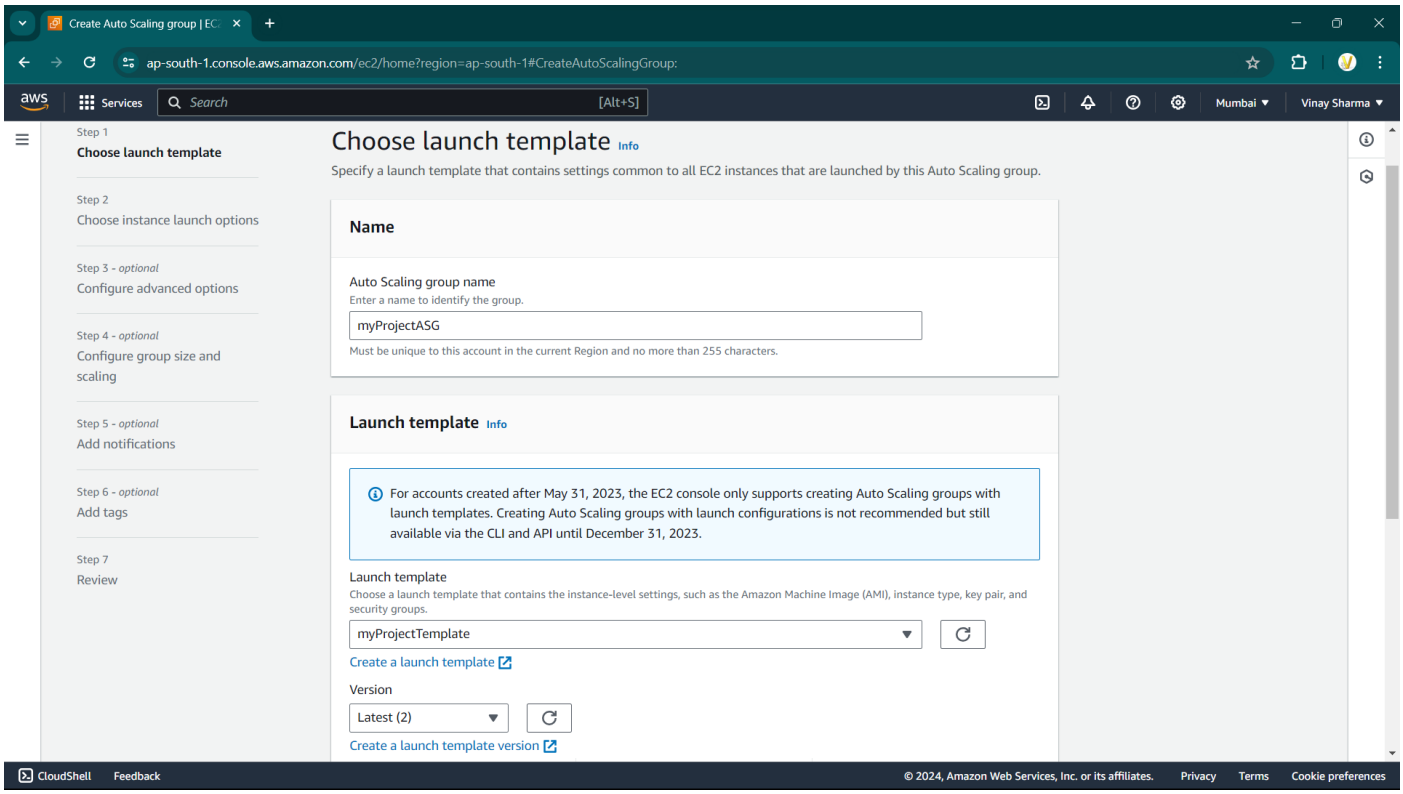

- *Here Select Launch Templates that we have created yet.*
- *Click on next.*

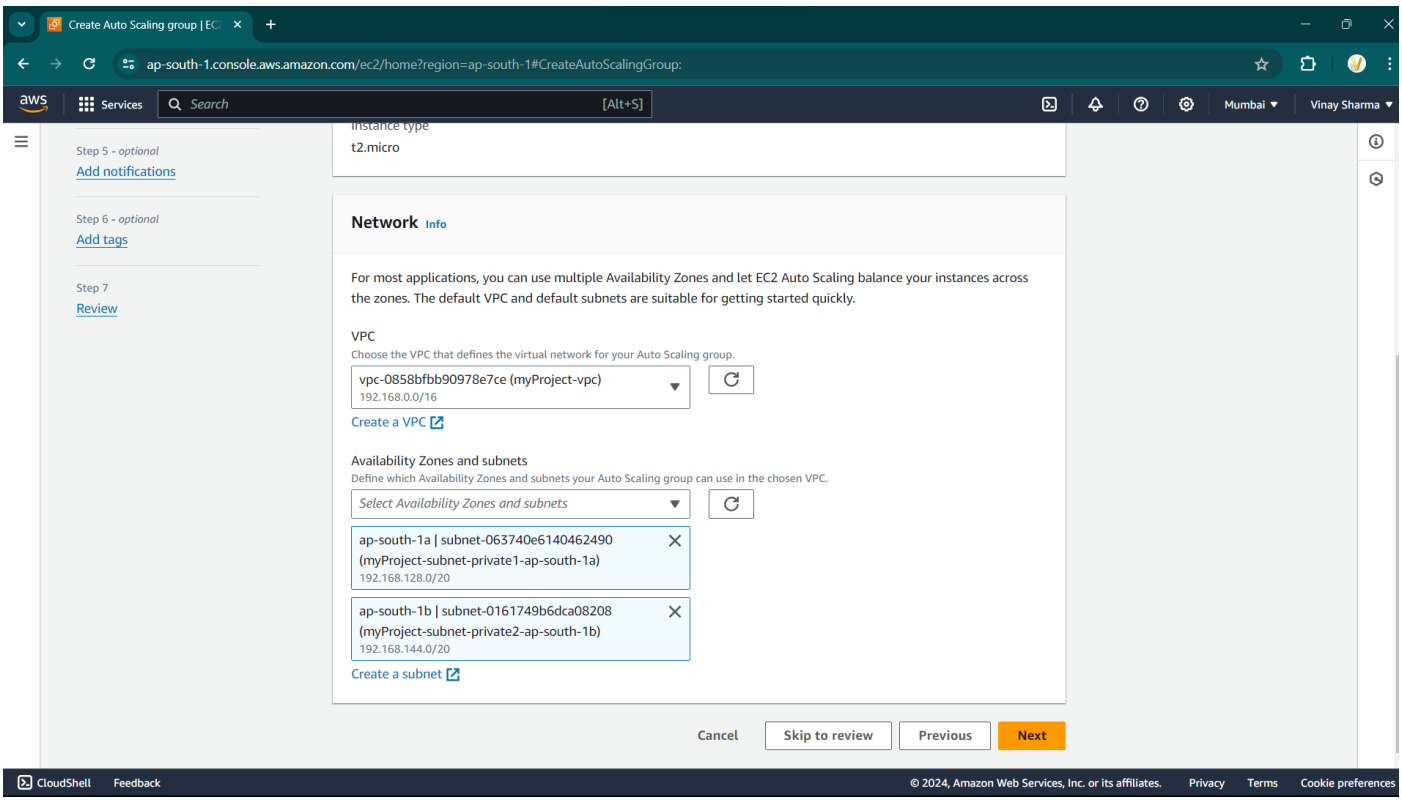

- *In VPC section, select our Custom VPC.*
- *In Subnets section select both of our private subnet, because we want our EC2 servers to run in private subnet.*

- *Click on next.*

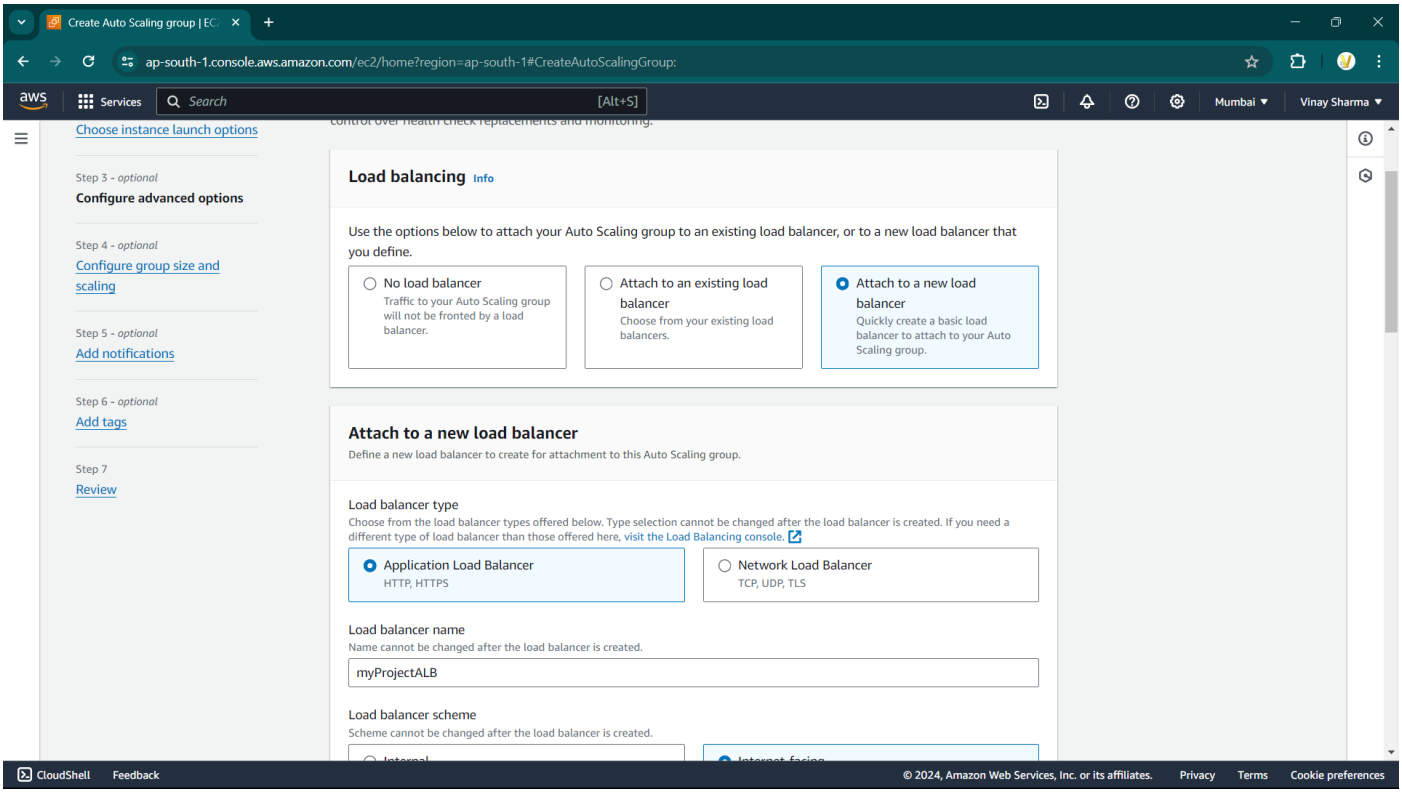

- *In load balancing option, choose attach new load balancer option; this will create and attach load balancer with our ASG.*

- *Choose Application load balancer as a type.*

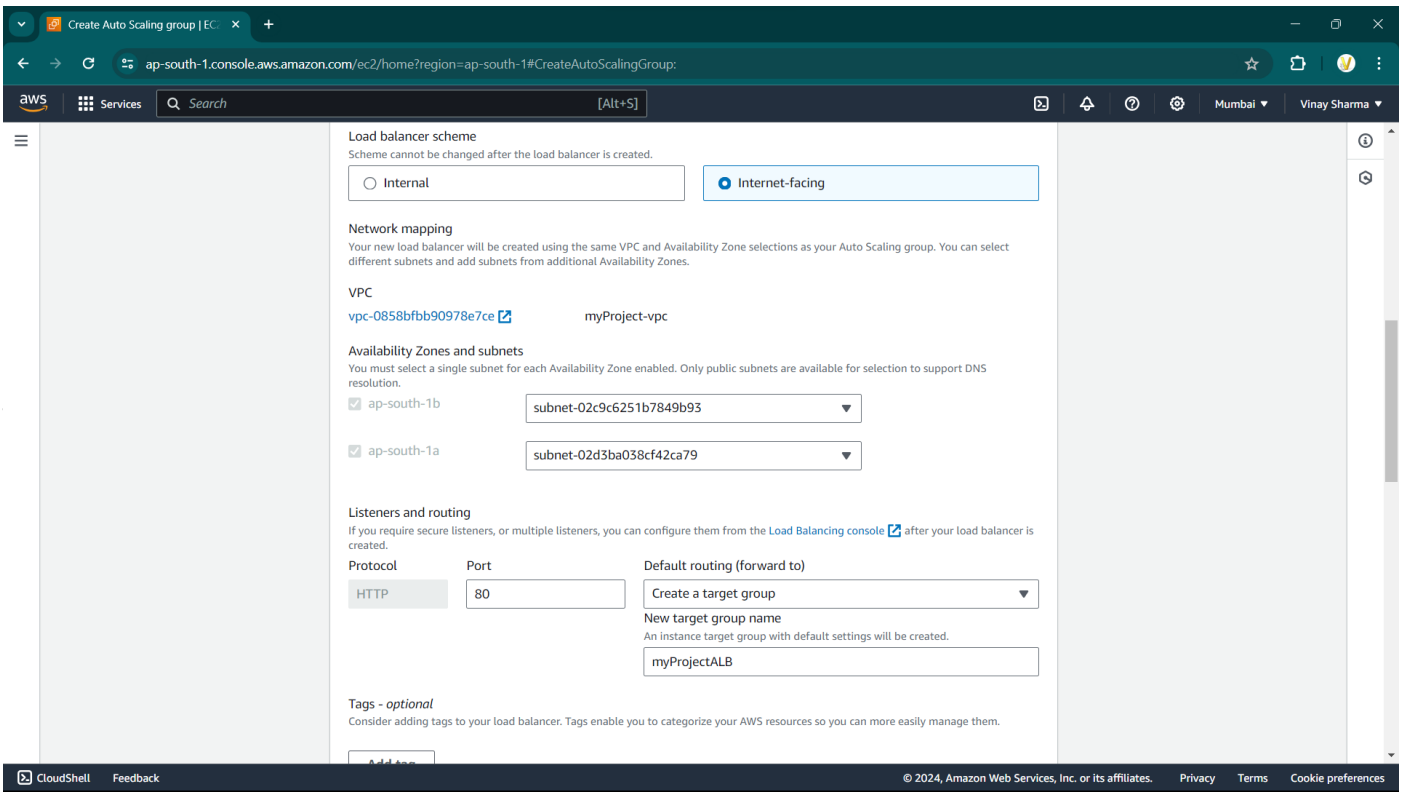

- *Make load balancer scheme as a internet-facing; this will make our load balancer to listen from internet.*

- *Choose public subnet; because we want to put our load balancer in public subnet.*

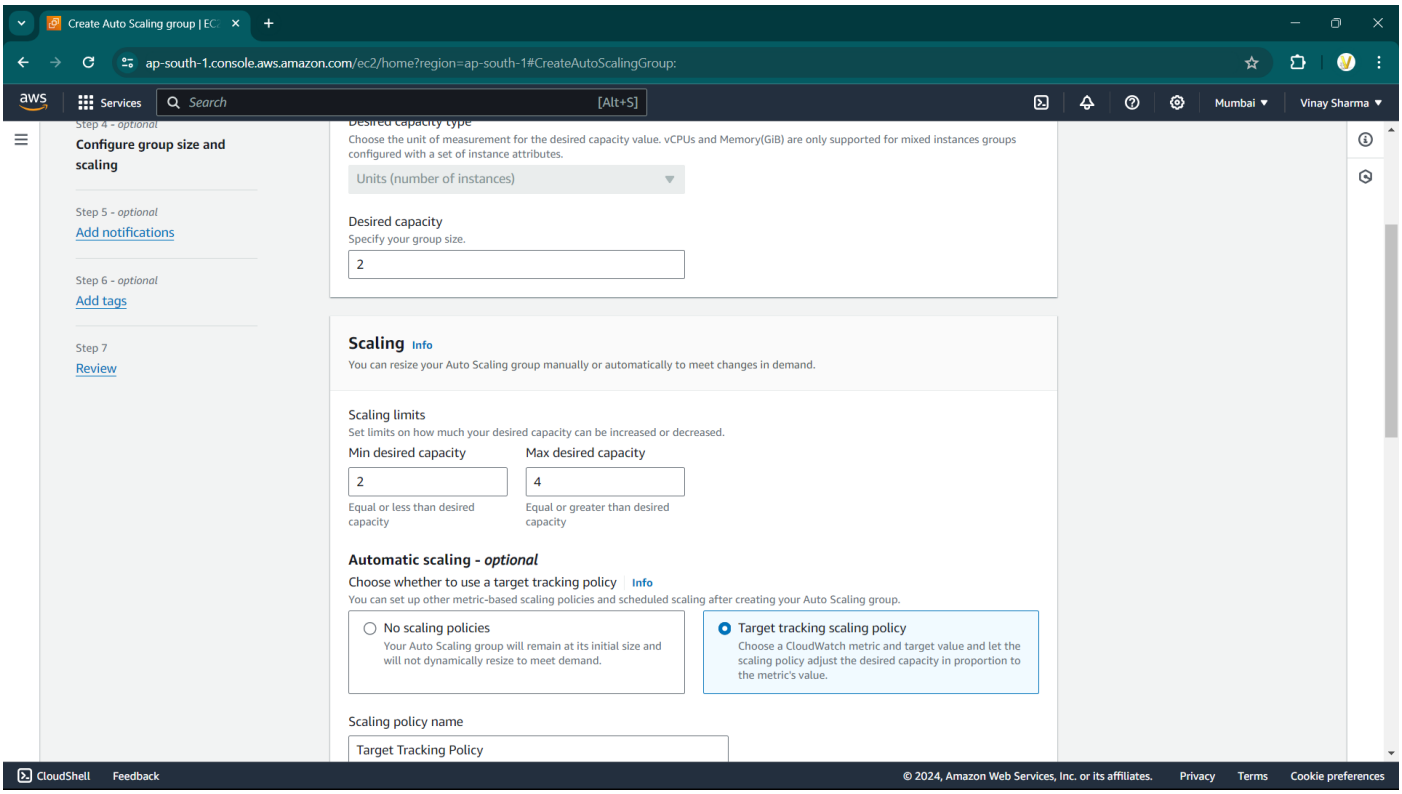

- *In a scaling policy, I want 2 instances each time, that's why I put 2 as a Min desired Capacity.*

- *In Max desired capacity I put 4; when load increases in our2 instances, so it will scale our number of instances to 4.*

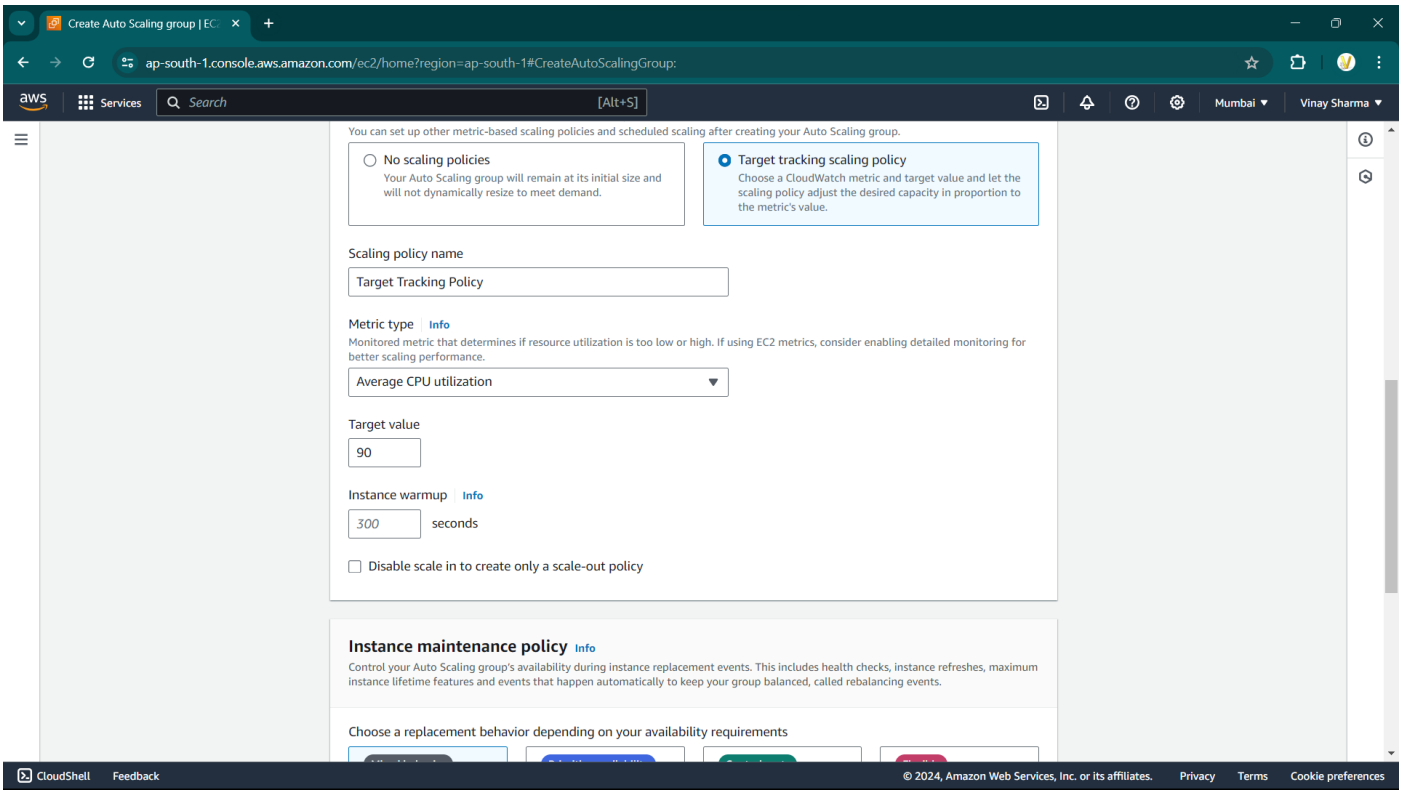

- *Here I choose metric type as an Average CPU Utilization and target value to 90; that means It will scale EC2 when our existing EC2 instances got more than 90% of CPU utilization.*

- *Now click on create Auto Scaling Group.*

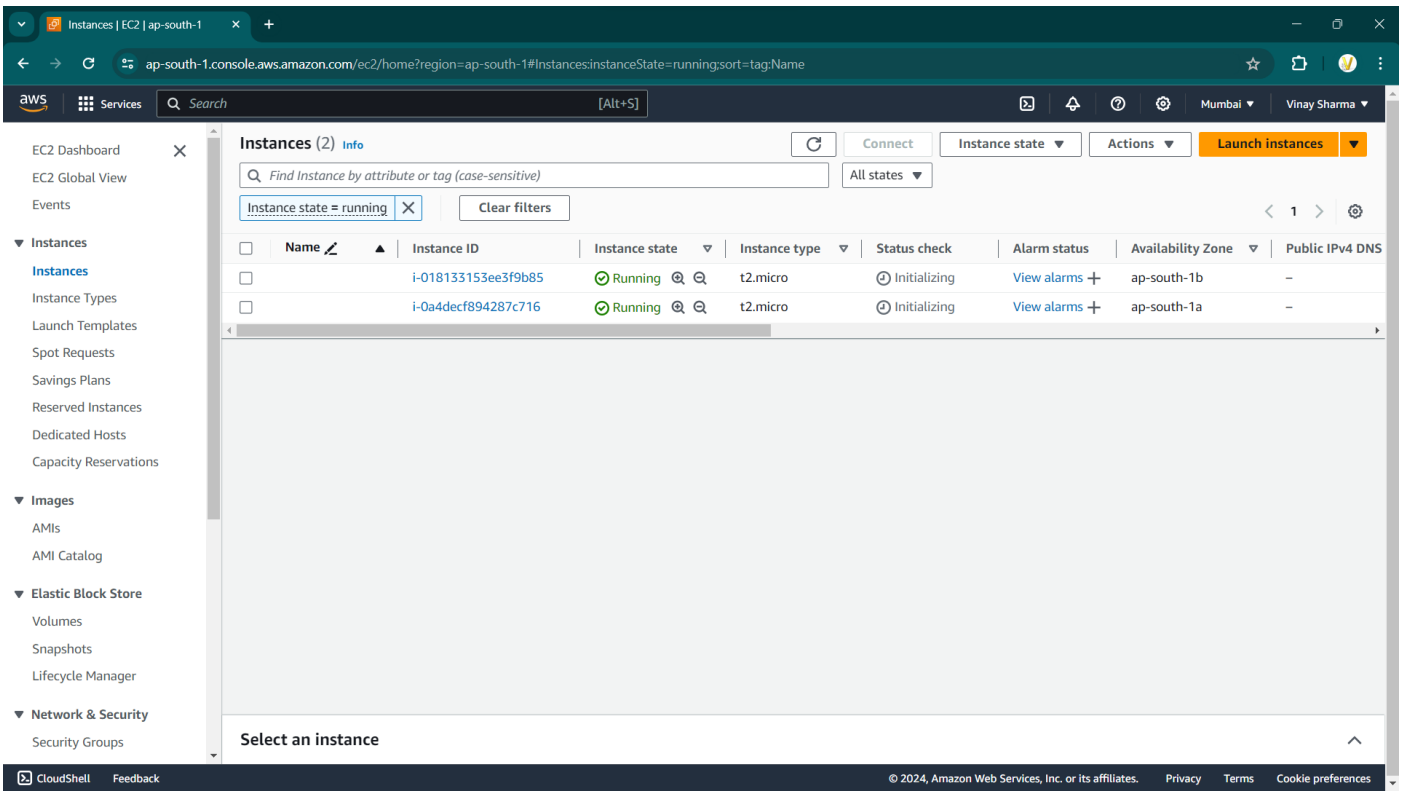

- *Here you can see that Auto Scaling Group has automatically created two instances for us; because we have defined desired capacity to 2.*

Bastion Host:

- *Here I have created one more instance that is known as bastion host; I have put this instance in public subnet of my VPC so that I am able to take SSH on it.*

- *I have transferred my index.html file (file that contains my website) & .pem file (that allows me to do SSH to my instances in private subnet) to bastion host using sFTP.*

Bastion host: A bastion host is a server whose purpose is to provide access to a private network from an external network, such as the Internet.

- *Here is the private IP address of both the machine in my private subnet:*

- *192.168.152.192 & 192.168.143.74*

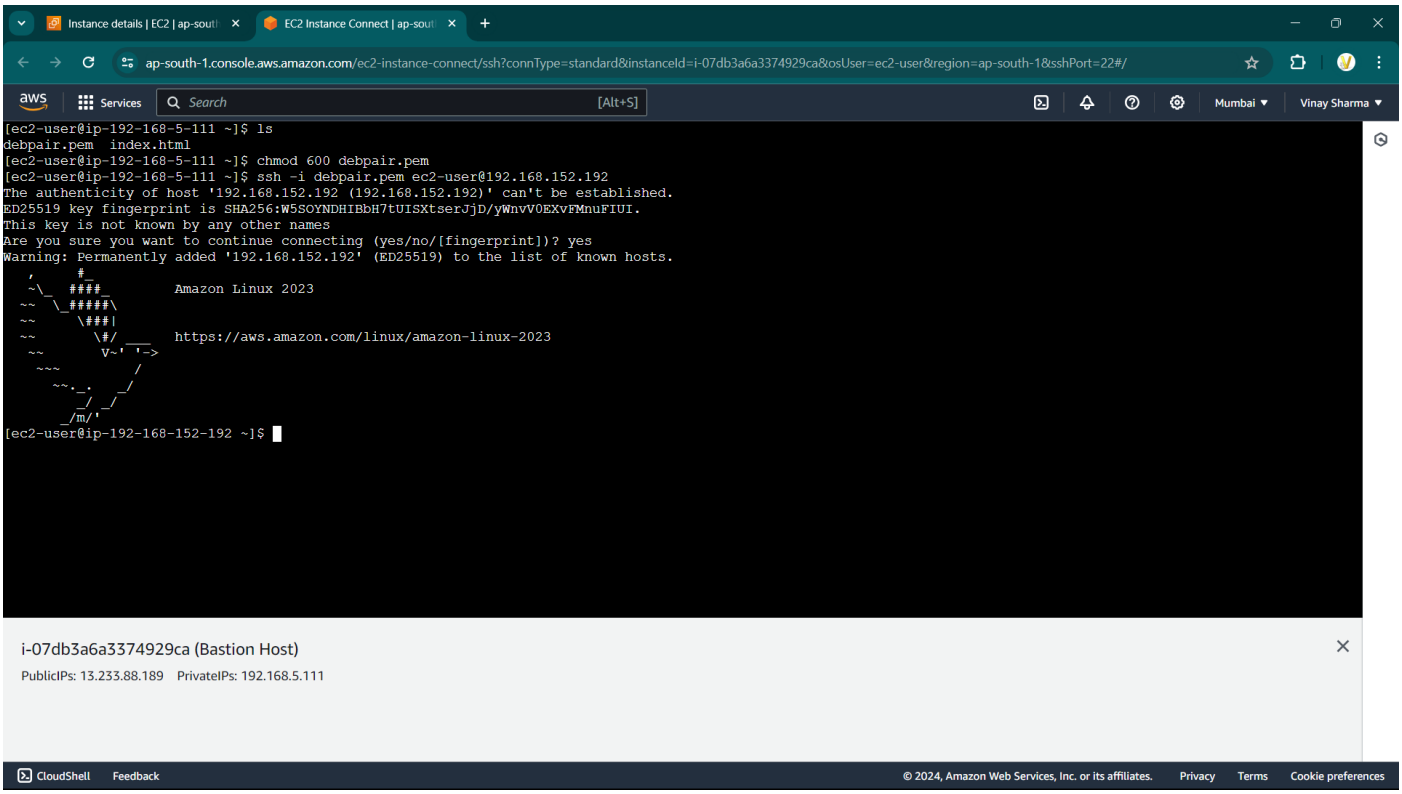

- *Now I have connected to my bastion host through EC2 instance connect.*

- *Then I uses SCP protocol to transfer index.html file to my both machines inside the private subnet.*

```
scp -i example.pem index.html ec2-user@192.168.152.192:/home/ec2-user
scp -i example.pem index.html ec2-user@192.168.143.74:/home/ec2-user
```

- *Then I uses SSH to connect both of the instances and copy index.html file from /home/ec2-user to /var/www/html/*

```
ssh -i example.pem ec2-user@192.168.152.192
sudo cp index.html /var/www/html
exit
ssh -i example.pem ec2-user@192.168.143.74
```

```
sudo cp index.html /var/www/html
exit
```

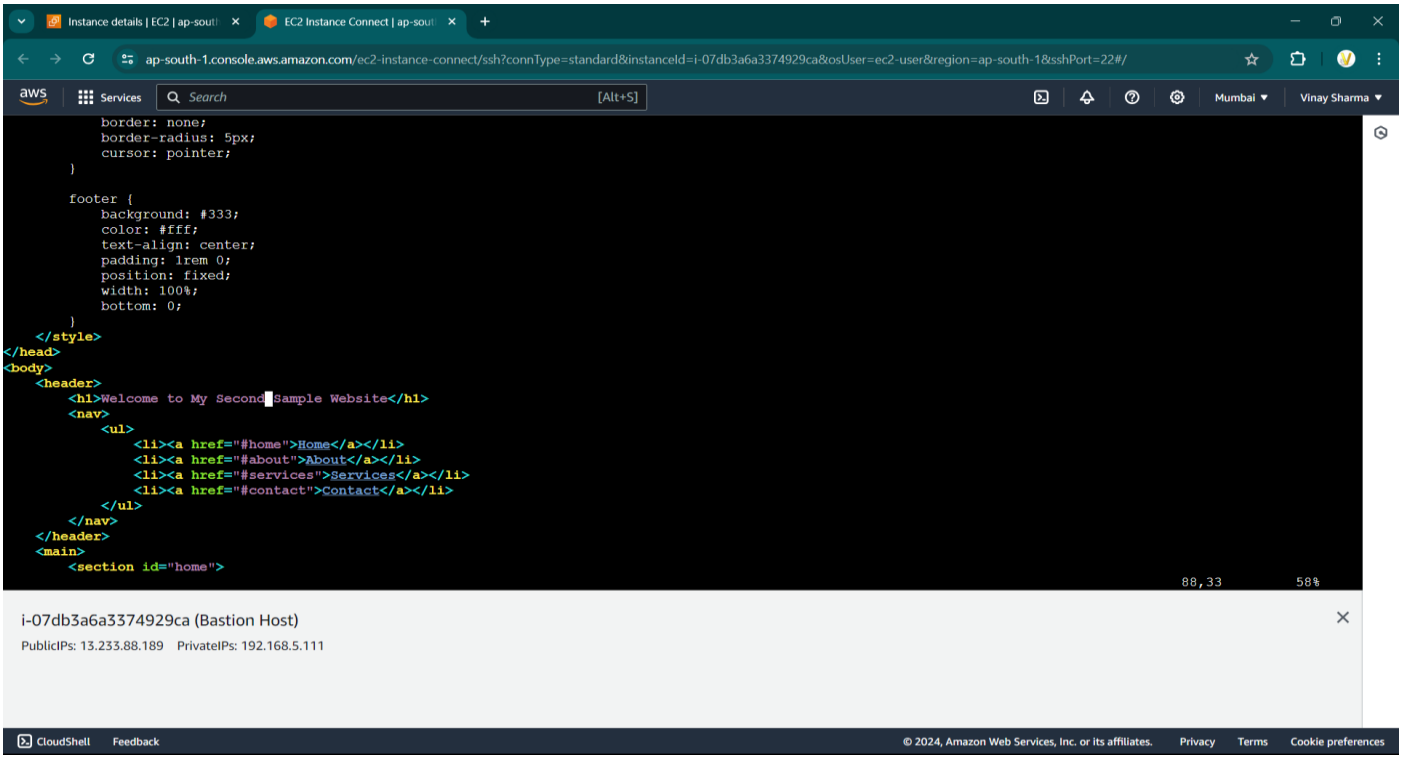

- *In one of the machine I have changed title of my website to ensures load balancing.*

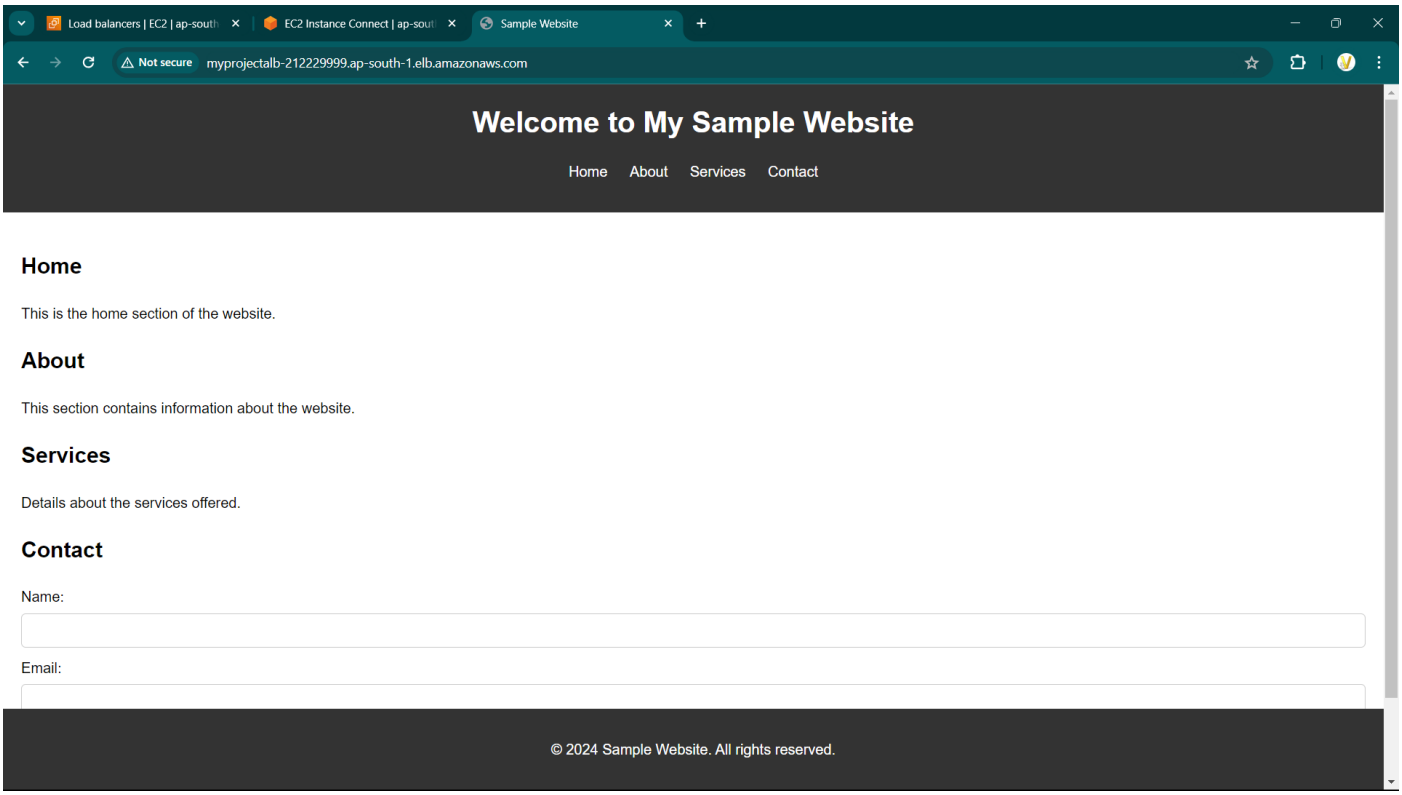

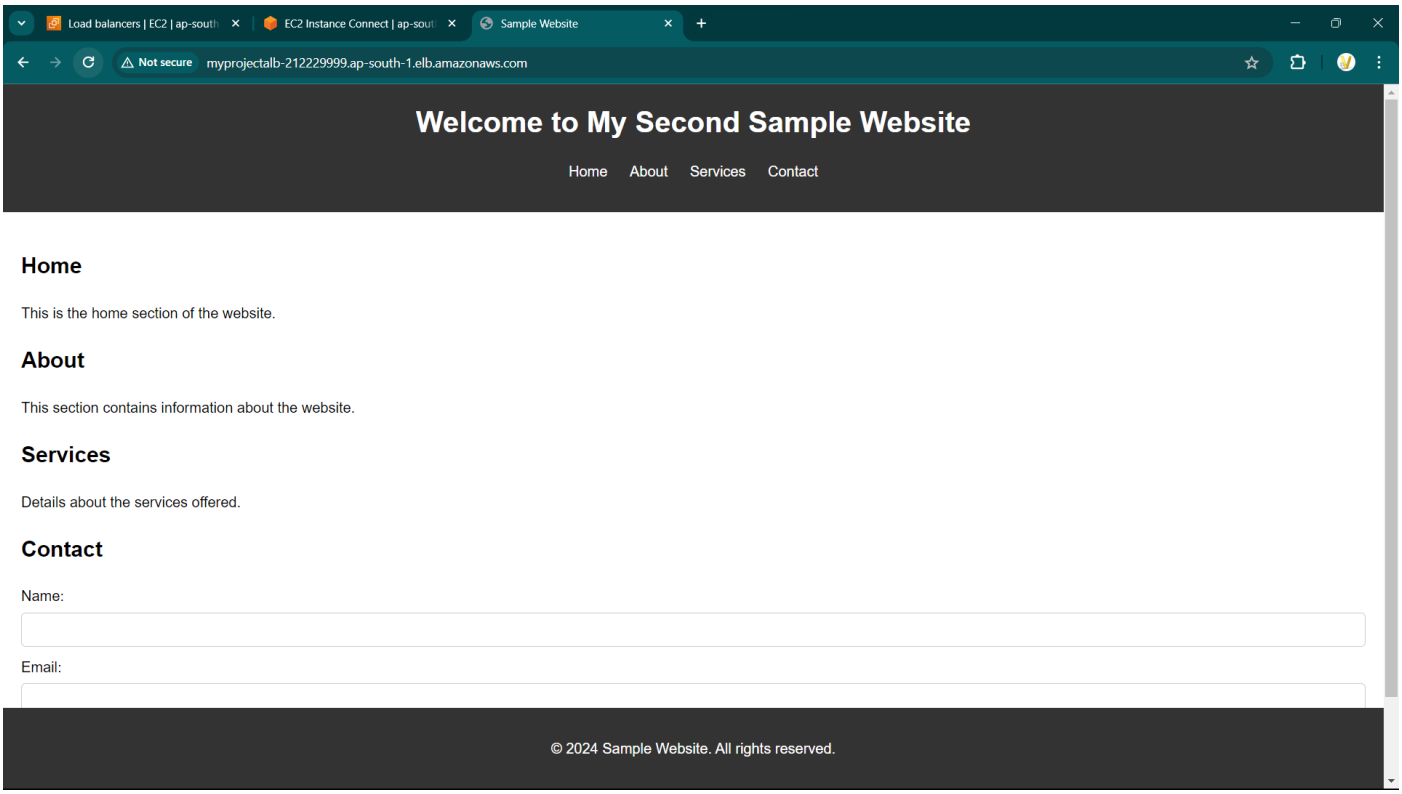

- *Using load balancer's DNS name, my website is got accessed. When I refresh my page, so you can see that my request is got redirected to another instance; that means my load balancer works fine!!!*

Let's Add Entry in DNS:

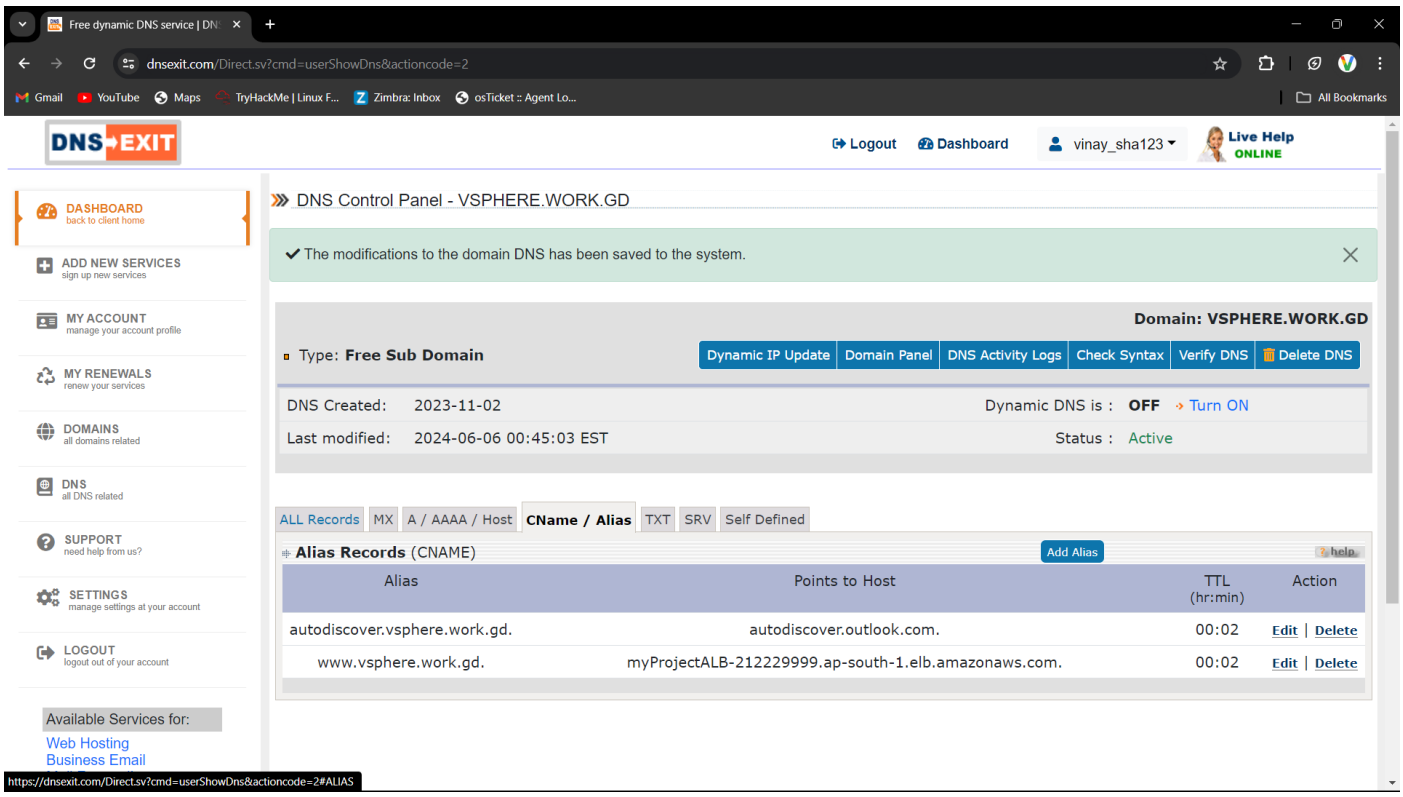

- *I want my website to get accessed using my own domain name "vsphere.work.gd". so I have created CNAME record that points www.vshpere.work.gd to my load balancer's DNS name.*

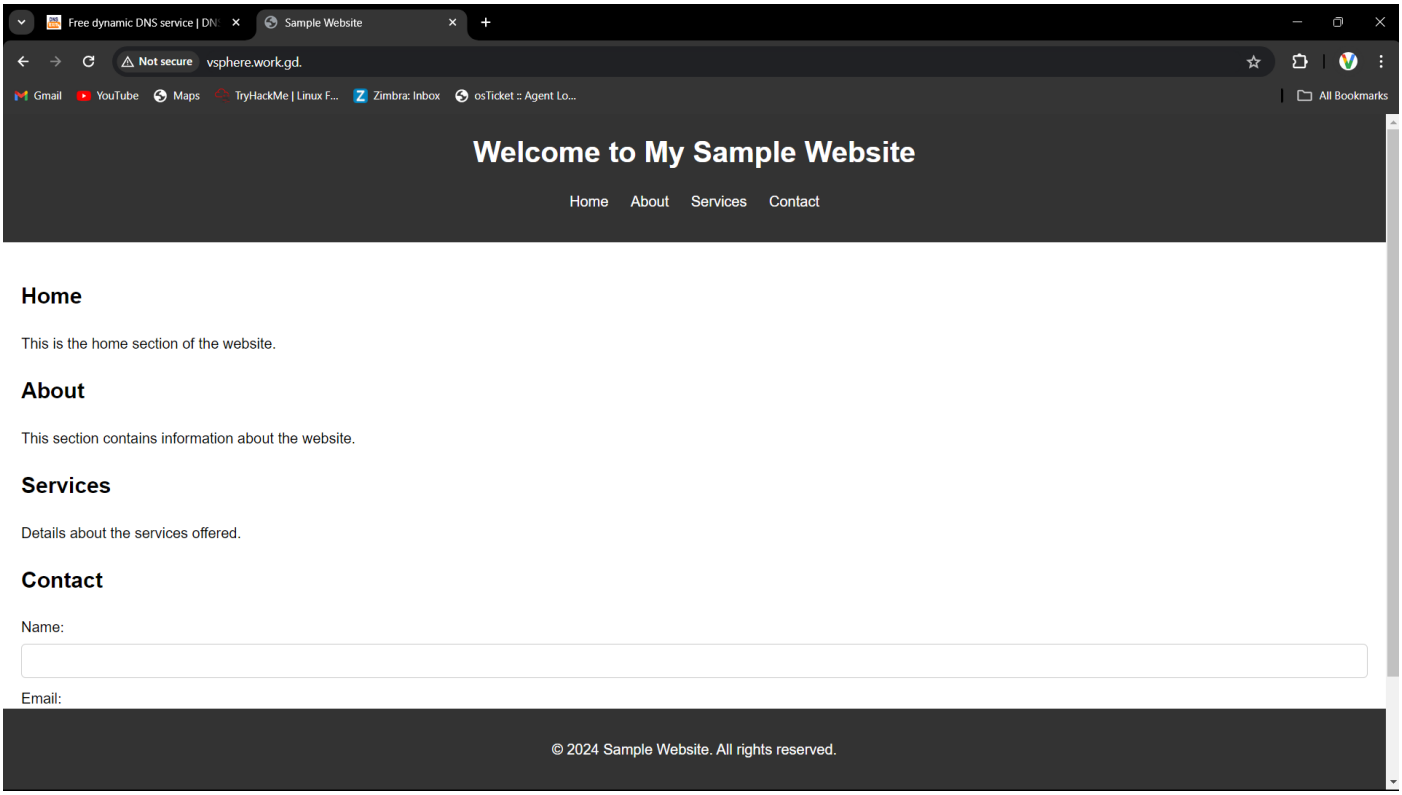

- *Now you can see that, my website is got accessed through my own domain name.*

Let's Generate SSL Certificate:

SSL Certificate: digital certificate that authenticates a website's identity and enables an encrypted connection.

Amazon Certificate Manager: AWS Certificate Manager (ACM) is a service that lets you easily provision, manage, and deploy public and private Secure Sockets Layer/Transport Layer Security (SSL/TLS) certificates.

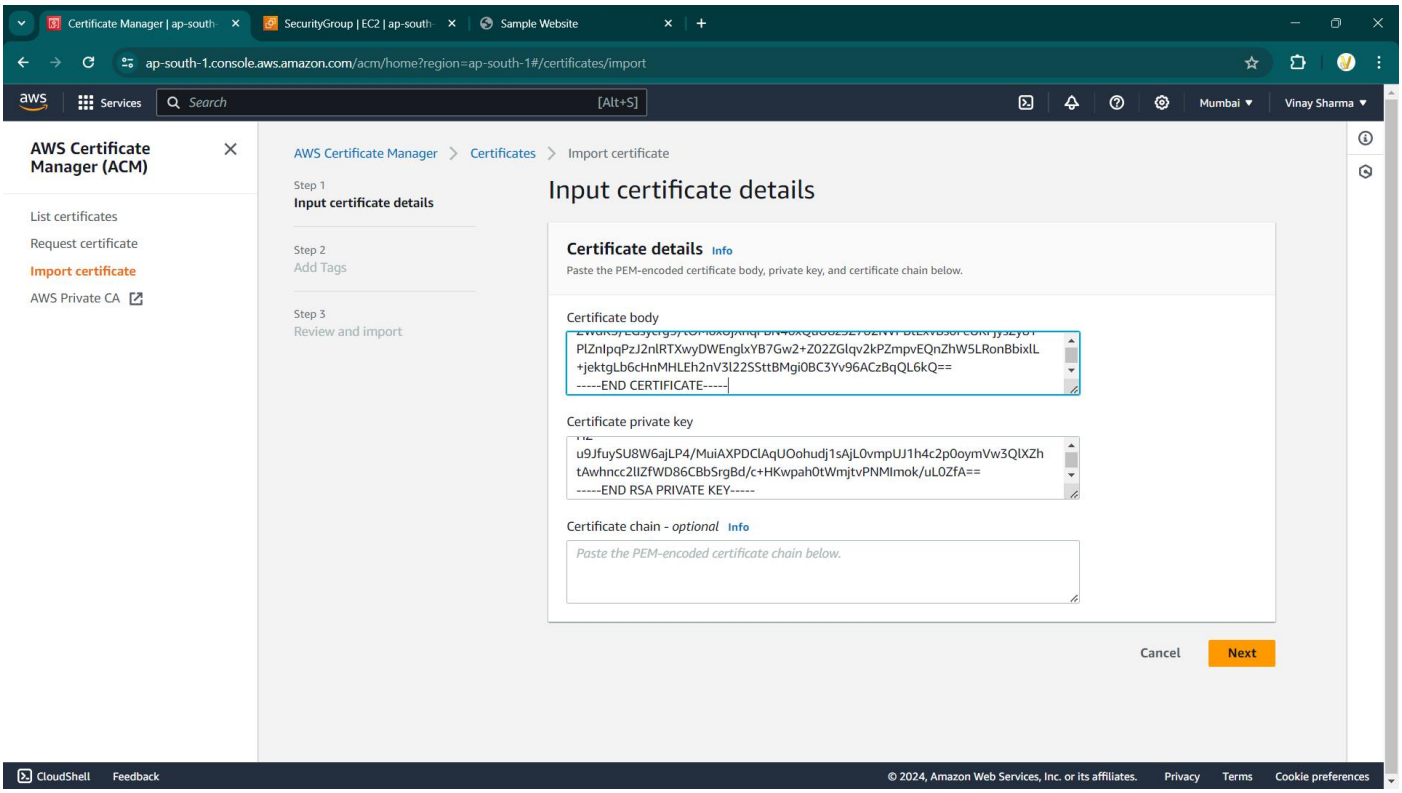

- *Currently my website is got accessed using http protocol, which is probably not secure.*

- *I have generated SSL certificate for my website from domain registrar, so that my website runs on HTTPS.*

- *I have provided .req file to my domain registrar and then i got certificate keys from it; that i have put in Amazon Certificate Manager.*

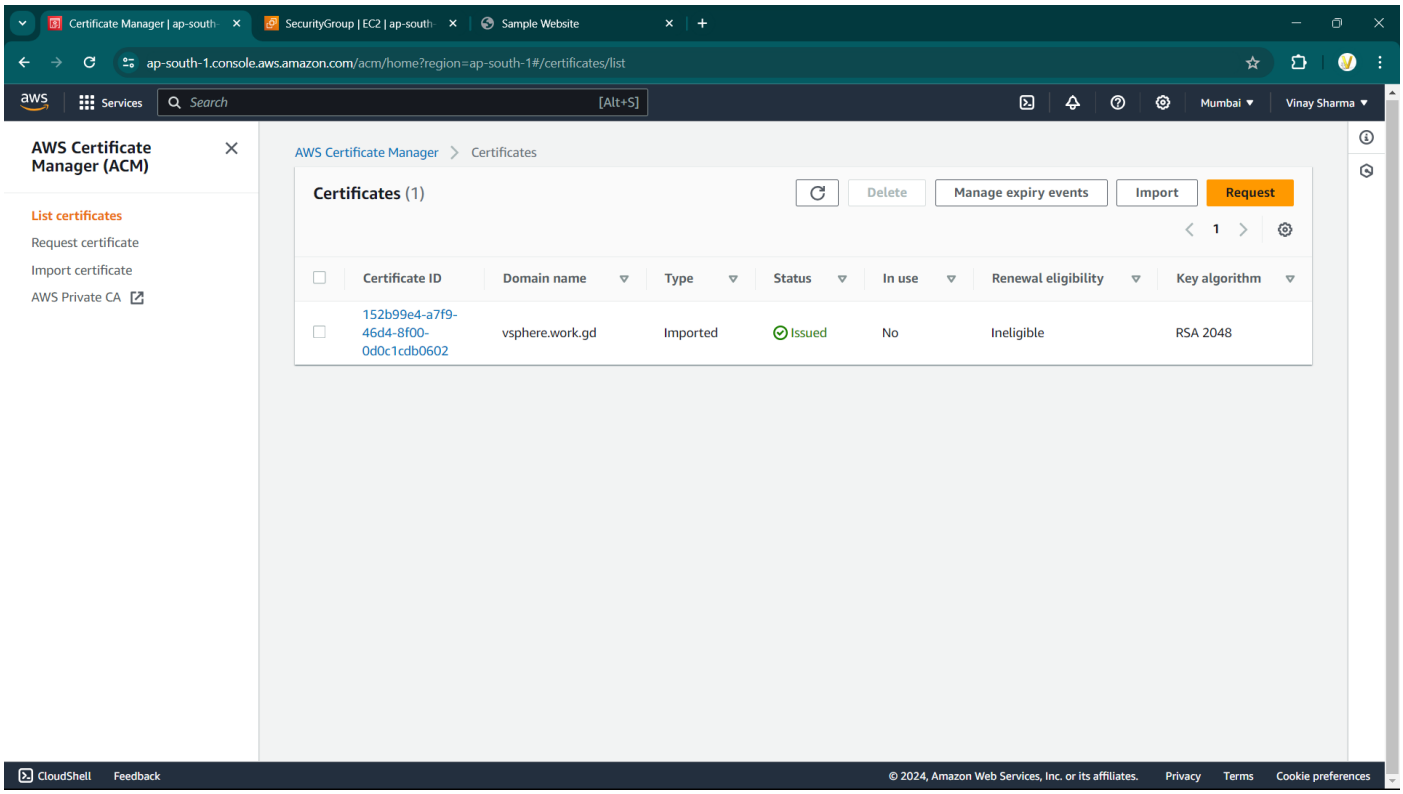

- *Here you can see, I have successfully imported my SSL certificate in ACM.*

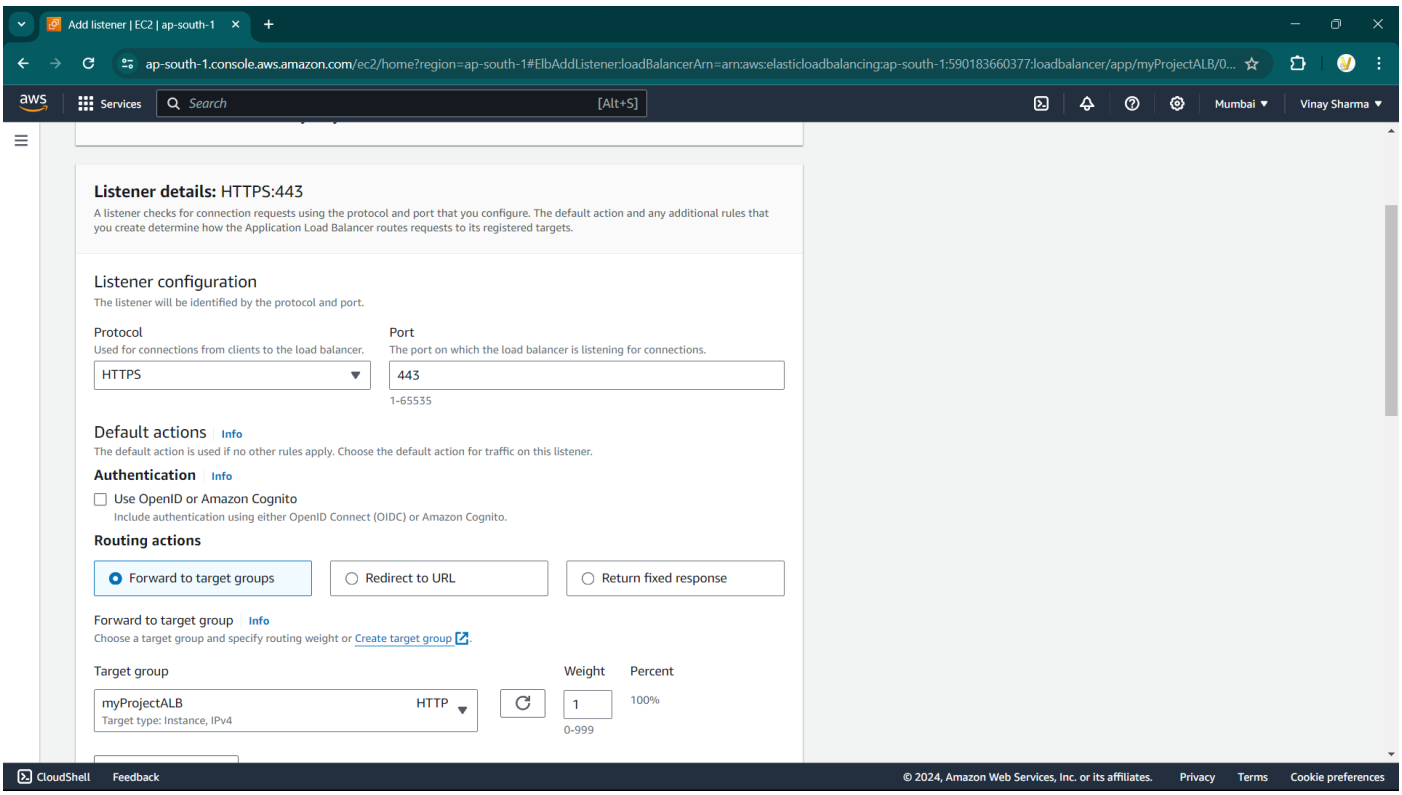

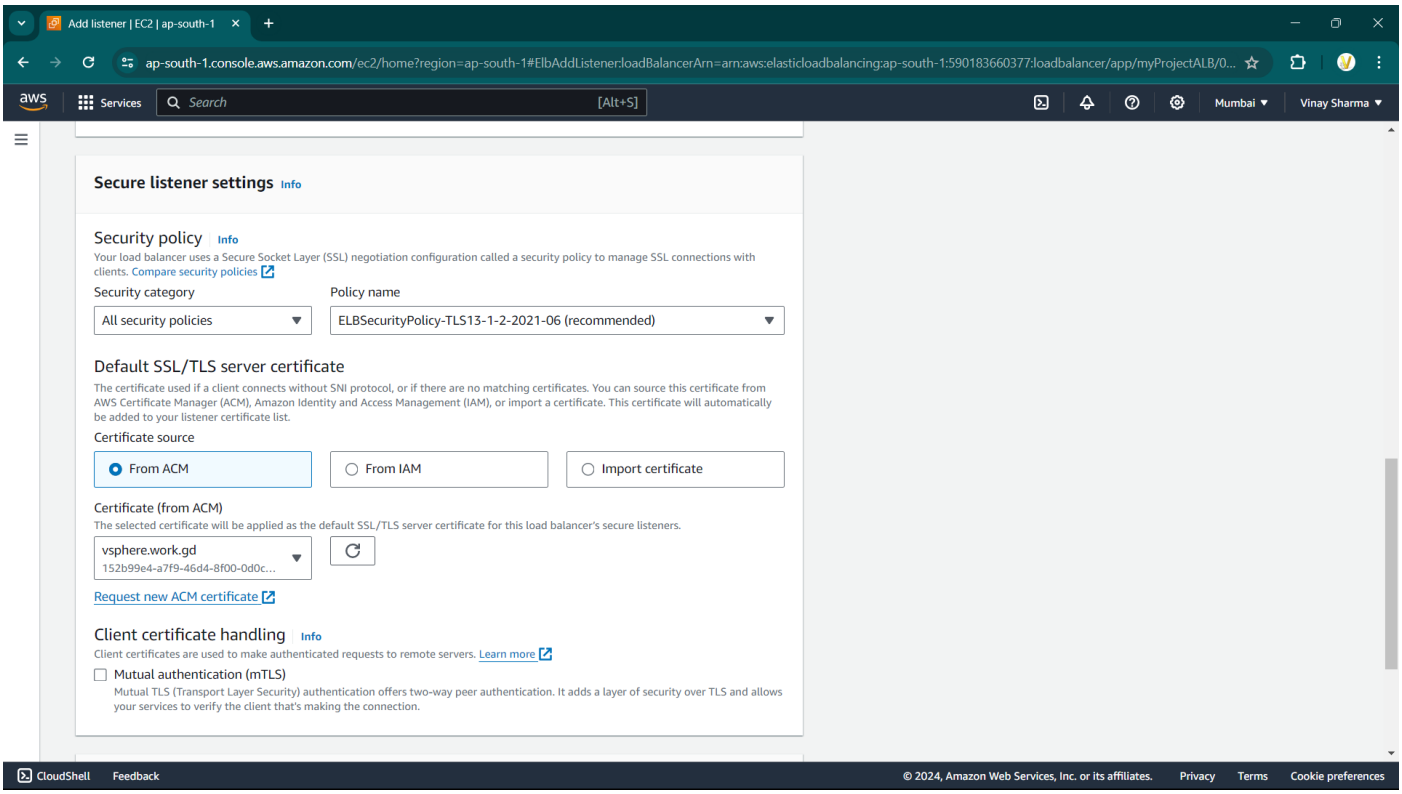

- *Now in load balancer's listeners rules, remove rule for HTTP and add new rule that works on HTTPS and in Default SSL/TLS certificate, use ACM as a certificate source and select our imported certificate and click on save.*

- *Now you can see that our website is got accessed through HTTPS protocol.*

Implementation

The implementation of the project involved several key steps to deploy a highly available and secure web application on AWS. Below is a detailed account of the implementation process:

1. Setting Up the AWS Environment

- **AWS Account**: Ensure you have an AWS account set up and ready to use.
- **IAM Roles**: Create necessary IAM roles and policies to manage access and permissions.

2. Creating a Custom VPC

- **VPC Creation**: Navigate to the VPC Dashboard and create a custom VPC with a suitable IP address range (e.g., 10.0.0.0/16).
- **Subnets**: Create subnets within the VPC:
  - **Public Subnet**: For resources that need internet access (e.g., NAT Gateway).
  - **Private Subnets**: For hosting the web servers to enhance security.

3. Configuring Security Groups and NACLs

- **Security Groups**: Define security groups to control inbound and outbound traffic for your instances.

    - Web Server Security Group: Allow traffic from the load balancer only.
    - Load Balancer Security Group: Allow HTTP/HTTPS traffic from the internet.
- **NACLs**: Configure Network Access Control Lists (NACLs) to add an additional layer of security at the subnet level.

4. Setting Up the Load Balancer

- **ELB Creation**: Create an Elastic Load Balancer (ELB) through the EC2 Dashboard.
- **Listeners and Target Groups**: Configure listeners for HTTP and HTTPS. Create target groups and register your web servers with these target groups.
- **SSL Certificate**: Obtain an SSL certificate from AWS Certificate Manager (ACM) and attach it to the load balancer for HTTPS traffic.

5. Deploying Web Servers

- **Launch EC2 Instances**: Launch EC2 instances in the private subnets to host your web application.
- **Web Server Configuration**: Install necessary software (e.g., Apache, Nginx) and deploy your web application on these instances.
- **Auto Scaling**: Set up an Auto Scaling Group to automatically adjust the number of instances based on demand.

6. Configuring Route Tables and Internet Gateway

- **Route Tables**: Create and associate route tables for each subnet. Ensure the private subnets have routes to the NAT Gateway for internet access.

- **Internet Gateway**: Attach an internet gateway to your VPC and configure the public subnet route table to allow internet access.

7. DNS and Domain Configuration

- **Domain Purchase**: Purchase a domain through AWS Route 53 or another domain registrar.
- **DNS Setup**: Create a hosted zone in Route 53 and configure DNS records to point to the load balancer.

8. Testing and Validation

- **Accessing the Application**: Test accessing the web application through the load balancer's public DNS name and ensure it routes traffic correctly.
- **Security Testing**: Perform security tests to ensure that the instances in private subnets are not directly accessible from the internet.
- **Load Testing**: Conduct load testing to verify that the application can handle expected traffic volumes.

By following these steps, the web application was successfully deployed on AWS with a robust architecture that ensures high availability, security, and scalability. The implementation process involved configuring various AWS services and components to create a secure and efficient environment for the application.

Conclusion

This project demonstrates the successful deployment of a highly available and secure web application on AWS. By leveraging an Elastic Load Balancer, multi-AZ deployment, and a custom VPC with private subnets, we ensured fault tolerance, scalability, and network isolation. Implementing security groups, NACLs, a custom domain, and an SSL certificate further enhanced the security and reliability of the application. Overall, this architecture provides a robust and secure environment, ensuring a seamless and secure experience for users.

References

- Youtube
- Udemy
- AWS Solutions Architect Study Guide v4